



Big Data Engineer Bootcamp

Code 1




Agenda

- **Project Structure**
- Dev Environment
- Work with Zookeeper
- Work with Kafka
- Work with Cassandra

Project Structure

- Apache Kafka
- Apache Zookeeper
- Apache Cassandra





Agenda

- Project Structure
- **Dev Environment**
- Work with Zookeeper
- Work with Kafka
- Work with Cassandra

Requirement

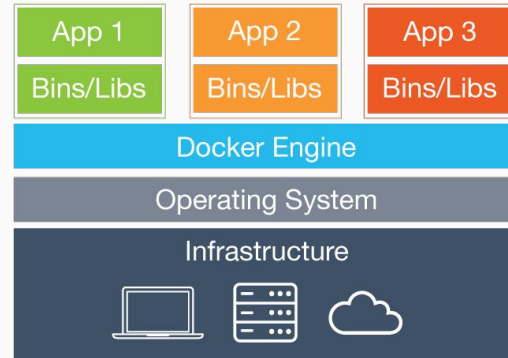
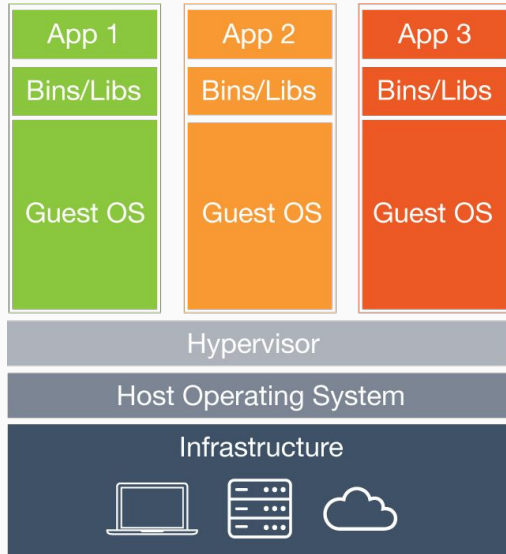
- Docker
- Virtualbox
- Docker-machine
 - 请一定要安装, 不然会导致环境不一致
- wget
- tar
- Scala
- SBT
- Python
- Pip

What is Docker

- A tool to package and deploy applications inside containers
 - Containers are isolated environments
- Developed by Solomon Hykes in Dotcloud
- Open-sourced in March 2013
- Grown into a platform
 - Docker Compose, Docker Swarm
 - Docker Image Hosting
 - Container Hosting



Why Docker



Docker Internal

- Client - Server Architecture
 - Docker command line = client 客户端可以连接任何远端的docker daemon
 - Docker daemon = server
- Use Linux functionalities (unavailable on MacOS or Windows)
 - cgroup
 - namespaces
 - Etc



Docker on Mac/Docker Machine

- 请一定要安装docker machine
- Docker on Mac默认会在本地启动一个小型的虚拟机, 会导致后面的命令代码不一致

Play with Docker

- `docker-machine create --driver virtualbox --virtualbox-cpu-count 2 --virtualbox-memory 2048 bigdata`
 - 创建一个名字叫做bigdata的虚拟机
 - 在这次课中, 请一定要用`docker machine`创建虚拟机再运行`docker`命令, 不然会导致命令无法正常执行
- `Docker-machine ip bigdata`
 - 这个命令可以帮助你看到bigdata这个虚拟机的ip地址
 - 后面的很多命令咱们都会用这个命令来获取正确的ip地址
- `eval $(docker-machine env bigdata)`
 - 注意这一个命令是必需的, 能够帮助你的`docker`客户端跟服务器通信, 每一个新的`terminal`窗口都需要输入这个命令
- `docker run -d -p 3000:3000 unclebarney/chit-chat`
- `docker images`
- `docker ps`



Docker on Windows/Docker Machine

- 安装Docker Toolbox (Include: VirtualBox, Docker-machine, Docker-compose)
 - Instruction: https://docs.docker.com/toolbox/toolbox_install_windows/
 - Download: <https://www.docker.com/products/docker-toolbox>
- 安装好后, 新开一个terminal, 然后使用docker-machine ls确认能否正确运行
 - 会列出已有的虚拟机

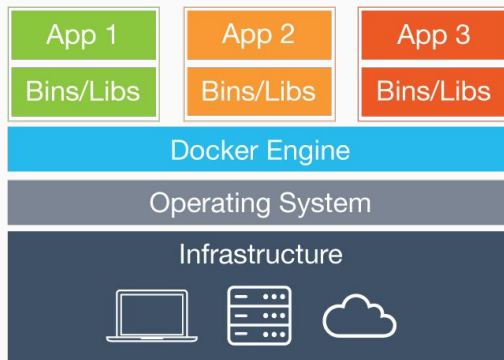
Play with Docker-Windows

- `docker-machine create --driver virtualbox --virtualbox-cpu-count 2 --virtualbox-memory 2048 bigdata`
 - 创建一个名字叫做bigdata的虚拟机
 - 在这次课中, 请一定要用`docker machine`创建虚拟机再运行`docker`命令, 不然会导致命令无法正常执行
- `Docker-machine ip bigdata`
 - 这个命令可以帮助你看到bigdata这个虚拟机的ip地址
 - 后面的很多命令咱们都会用这个命令来获取正确的ip地址
- `docker-machine env --shell cmd default`
- `FOR /f "tokens=*" %i IN ('docker-machine env --shell cmd bigdata') DO %i`
 - 注意这个命令是必需的, 能够帮助你的`docker`客户端跟服务器通信, 每一个新的`terminal`窗口都需要输入这个命令
- `docker run -d -p 3000:3000 unclebarney/chit-chat`
- `docker images`
- `docker ps`



Dev Environment

- Run all the servers/components as docker container
- For example
 - App 1 = zookeeper
 - App 2 = kafka
 - App 3 = cassandra
- Allow fast iteration





Agenda

- Project Structure
- Dev Environment
- **Work with Zookeeper**
- Work with Kafka
- Work with Cassandra

Start Zookeeper Server

- `docker run -d -p 2181:2181 -p 2888:2888 -p 3888:3888 --name zookeeper confluent/zookeeper`
- `docker images`
- `docker ps`



Get Zookeeper CLI

- Download using shell commands (MacOS, Linux, Unix)
 - `wget http://apache.mirrors.ionfish.org/zookeeper/zookeeper-3.4.8/zookeeper-3.4.8.tar.gz`
 - `tar xvf zookeeper-3.4.8.tar.gz`
 - `mv zookeeper-3.4.8 zookeeper`
 - `rm zookeeper-3.4.8.tar.gz`
- Download directly (Windows)
 - `http://www.apache.org/dyn/closer.cgi/zookeeper/`

- `cd zookeeper/bin` (MacOS, Linux, Unix)
- `./zkCli.sh -server `docker-machine ip bigdata`:2181`
- `./zkCli.sh -server localhost:2181`
- `cd zookeeper/bin` (Windows)
- `docker-machine ip bigdata`
 - memorize virtual machine ip (ex. 192.168.99.100)
- `zkCli.cmd -server 192.168.99.100:2181`



Browse Znode Data

- `ls /`
- `ls /zookeeper`
- `get /zookeeper/quota`



Create Znode Data

- `create /workers "bittiger"`
- `ls /`
- `ls /workers`
- `get /workers`

```
bittiger
cZxid = 0x2
ctime = Sat Aug 20 22:27:29 PDT 2016
mZxid = 0x2
mtime = Sat Aug 20 22:27:29 PDT 2016
pZxid = 0x2
cversion = 0
dataVersion = 0
aclVersion = 0
ephemeralOwner = 0x0
dataLength = 8
numChildren = 0
```



Delete Znode Data

- `delete /workers`
- `ls /`
- `ls /workers`
- `get /workers`



Create Ephemeral Znode Data

- `create -e /workers "unclebarney"`
- `ls /`
- `ls /workers`
- `get /workers`

```
unclebarney
cZxid = 0x9
ctime = Sat Aug 20 22:34:44 PDT 2016
mZxid = 0x9
mtime = Sat Aug 20 22:34:44 PDT 2016
pZxid = 0x9
cversion = 0
dataVersion = 0
aclVersion = 0
ephemeralOwner = 0x156ab8464e60002
dataLength = 11
numChildren = 0
```



Watcher

- `get /workers true`





Agenda

- Project Structure
- Dev Environment
- Work with Zookeeper
- **Work with Kafka**
- Work with Cassandra

Dependencies

- `scala -version`
- `sbt --version`
- `python --version`
- `pip --version`



Start Kafka Server

- `docker run -d -p 9092:9092 -e KAFKA_ADVERTISED_HOST_NAME=`docker-machine ip bigdata` -e KAFKA_ADVERTISED_PORT=9092 --name kafka --link zookeeper:zookeeper confluent/kafka`
- `docker images`
- `docker ps`



Get Kafka CLI

- Download using command line (MacOS, Unix, Linux)
 - `wget http://apache.mirrors.ionfish.org/kafka/0.10.0.1/kafka_2.11-0.10.0.1.tgz`
 - `tar xvf kafka_2.11-0.10.0.1.tgz`
 - `mv kafka_2.11-0.10.0.1 kafka`
 - `rm kafka_2.11-0.10.0.1.tgz`
- Download directly (Windows)
 - `https://www.apache.org/dyn/closer.cgi?path=/kafka/0.10.0.1/kafka_2.11-0.10.0.1.tgz`



Create Kafka Topic

- (MacOS, Unix, Linux)
- `./kafka-topics.sh --create --zookeeper `docker-machine ip bigdata` --replication-factor 1 --partitions 1 --topic bigdata`
- `./kafka-topics.sh --list --zookeeper `docker-machine ip bigdata``
- (Windows)
- `docker-machine ip bigdata`
 - memorize virtual machine ip (ex. **192.168.99.100**, please change accordingly. I will use this in all following slices for windows)
- `./windows/kafka-topics.bat --create --zookeeper `docker-machine ip bigdata` --replication-factor 1 --partitions 1 --topic bigdata`
 - problem of “java.lang.classnotfoundException” please look up classpath in environ var
- `./windows/kafka-topics.bat --list --zookeeper `docker-machine ip bigdata``



Look up on Zookeeper

- (MacOS, Unix, Linux)
- `./zkCli.sh -server `docker-machine ip bigdata`:2181`
- `ls /`

- (Windows)
- `zkCli.cmd -server 192.168.99.100:2181`
- `ls /`



Produce Messages

- (MacOS, Unix, Linux)
- `./kafka-console-producer.sh --broker-list `docker-machine ip bigdata`:9092 --topic bigdata`
- (Windows)
- `kafka-console-producer.bat --broker-list 192.168.99.100:9092 --topic bigdata`



Consume Messages

- (MacOS, Unix, Linux)
 - `./kafka-console-consumer.sh --zookeeper `docker-machine ip bigdata`:2181 --topic bigdata`
 - `./kafka-console-consumer.sh --zookeeper `docker-machine ip bigdata`:2181 --topic bigdata --from-beginning`
-
- (Windows)
 - `kafka-console-consumer.bat --zookeeper 192.168.99.100:9092:2181 --topic bigdata`
 - `kafka-console-consumer.bat --zookeeper 192.168.99.100:9092:2181 --topic bigdata --from-beginning`



Look Into Kafka Broker

- `docker exec -it kafka bash`
- `cd /var/lib/kafka`
- `ls`





Agenda

- Project Structure
- Dev Environment
- Work with Zookeeper
- Work with Kafka
- **Work with Cassandra**

Start Cassandra Server

- `docker run -d -p 7199:7199 -p 9042:9042 -p 9160:9160 -p 7001:7001 --name cassandra cassandra:3.7`
- `docker images`
- `docker ps`



Get Cassandra CLI

- Download using commandline (Mac, Linux, Unix)
 - `wget http://apache.mirrors.ionfish.org/cassandra/3.7/apache-cassandra-3.7-bin.tar.gz`
 - `tar xvf apache-cassandra-3.7-bin.tar.gz`
 - `mv apache-cassandra-3.7 cassandra`
 - `rm apache-cassandra-3.7-bin.tar.gz`
- Direct Download(windows)
 - `http://www.apache.org/dyn/closer.lua/cassandra/3.7/apache-cassandra-3.7-bin.tar.gz`
- `./cqlsh `docker-machine ip bigdata` 9042`(Mac, Linux, Unix)
- `cqlsh.bat 192.168.99.100 9042`(windows)



Create Keyspace'

- `./cqlsh `docker-machine ip bigdata` 9042`
- `CREATE KEYSPACE "stock" WITH replication = {'class': 'SimpleStrategy', 'replication_factor': 1} AND durable_writes = 'true';`
- `USE stock;`
- `DESCRIBE KEYSPACE;`



Create Table

- `./cqlsh `docker-machine ip bigdata` 9042`
- `CREATE TABLE user (first_name text, last_name text, PRIMARY KEY (first_name));`
- `DESCRIBE TABLE user;`



Insert Data

- `./cqlsh `docker-machine ip bigdata` 9042`
- `INSERT INTO user (first_name, last_name) VALUES ('uncle', 'barney');`



Query Data

- `./cqlsh `docker-machine ip bigdata` 9042`
- `SELECT COUNT (*) FROM USER;`
- `SELECT * FROM user WHERE first_name='uncle';`
- `SELECT * FROM user WHERE last_name='barney';`



Look Into Cassandra Node

- `docker exec -it cassandra bash`
- `cd /var/lib/cassandra`
- `ls`



Delete Data

- `./cqlsh `docker-machine ip bigdata` 9042`
- `DELETE last_name FROM user WHERE first_name='uncle';`
- `DELETE FROM user WHERE first_name='uncle';`



Remove Table

- `./cqlsh `docker-machine ip bigdata` 9042`
- `TRUNCATE user;`
- `DROP TABLE user;`

