



Big Data Engineer Bootcamp

Part 1



Agenda

- More on Big Data
- Introduction to Kafka
- Introduction to Zookeeper
- Introduction to Cassandra
- Q&A



Agenda

- **More on Big Data**
- Introduction to Kafka
- Introduction to Zookeeper
- Introduction to Cassandra
- Q&A



More on Big Data

- **History of Big Data**
- Frameworks and Tools

History of Big Data

- 1991 - internet was born (yay)
- 1996 - digital storage is so cheap, cheaper than paper
- 1998 - Google was founded (yay)
- 1999 - the word Big Data first used in a paper
- 2001 - idea of 3V of data described by Doug Laney



Web 1.0

- Front End
 - HTML
 - CSS
 - JavaScript
- Dynamic Content
 - PHP
 - JSP
 - ASP.NET
 - Ruby
 - Perl
- Database
 - RDBMS
- Site Owner Generates Content



 [Auctions](#)  [Messenger](#)  [Check Email](#)   [What's New](#)  [Personalize](#)  [Help](#)

Personal Mail
you@your-domain.com

 Know when friends are online!
Click to download Yahoo! Messenger

NEW! Play ball!
free **Fantasy Baseball**

Search [advanced search](#)

Y! Shopping Depts: [Books](#), [CDs](#), [Computers](#), [DVDs](#) Stores: and [more](#)

Shop [Auctions](#) · [Classifieds](#) · [PayDirect](#) · [Shopping](#) · [Travel](#) · [Yellow Pgs](#) · [Maps](#) Media [Finance/Quotes](#) · [News](#) · [Sports](#) · [Weather](#)
Connect [Chat](#) · [Clubs](#) · [Experts](#) · [GeoCities](#) · [Greetings](#) · [Invites](#) · [Mail](#) · [Members](#) · [Messenger](#) · [Mobile](#) · [Personals](#) · [People Search](#)
Personal [Addr Book](#) · [Briefcase](#) · [Calendar](#) · [My Yahoo!](#) · [Photos](#) Fun [Games](#) · [Kids](#) · [Movies](#) · [Music](#) · [Radio](#) · [TV](#) [more...](#)

Yahoo! Auctions

- Bid, buy, or sell anything!

Categories

· [Antiques](#)
· [Cameras](#)
· [Coins](#)
· [Comic Books](#)

· [Computers](#)
· [Electronics](#)
· [Sports Cards](#)
· [Stamps](#)

Items

· [Longaberger](#)
· [PlayStation 2](#)
· [MP3 Players](#)
· [Golf Clubs](#)
· [Scooters](#)
· [Dale Earnhardt](#)
· [States Quarters](#)
· [Palm Pilots](#)

Got Something to Sell? [Auction it Now!](#)

Arts & Humanities

[Literature](#), [Photography](#)...

Business & Economy

[B2B](#), [Finance](#), [Shopping](#), [Jobs](#)...

Computers & Internet

[Internet](#), [WWW](#), [Software](#), [Games](#)...

Education

[College and University](#), [K-12](#)...

Entertainment

[Cool Links](#), [Movies](#), [Humor](#), [Music](#)...

Government

[Elections](#), [Military](#), [Law](#), [Taxes](#)...

Health

[Medicine](#), [Diseases](#), [Drugs](#), [Fitness](#)...

News & Media

[Full Coverage](#), [Newspapers](#), [TV](#)...

Recreation & Sports

[Sports](#), [Travel](#), [Autos](#), [Outdoors](#)...

Reference

[Libraries](#), [Dictionaries](#), [Quotations](#)...

Regional

[Countries](#), [Regions](#), [US States](#)...

Science

[Animals](#), [Astronomy](#), [Engineering](#)...

Social Science

[Archaeology](#), [Economics](#), [Languages](#)...

Society & Culture

[People](#), [Environment](#), [Religion](#)...

In the News

- [6.8 earthquake shakes Pacific Northwest; some injuries reported](#)
 - [Irish confirm spread of foot-and-mouth disease](#)
 - [FBI arrests 7 members of anti-Iranian terror group](#)
 - [NASA ends asteroid mission](#)
- [more...](#)

Marketplace

- [Tax Center](#) - forms, tips, online filing and more
- [Get your own Web domain](#)
- [Y! Careers](#) - find a new job!
- [Insurance](#) - [Auto](#), [Life](#), [Health](#), [Home](#) - get quotes, tips, more

Broadcast Events

- 8pm ET : [Heat vs. 76ers](#)
 - 8pm : Florida vs. [Vanderbilt](#)
 - 9pm : [NC State](#) vs. [North Carolina](#)
- [more...](#)

Inside Yahoo!

- **new!** Play free [Fantasy Baseball](#)
- [Astrology](#) - what's your sign?
- [Yahooligans!](#) - for kids
- [Oscar Pick'em](#) at Y! Movies

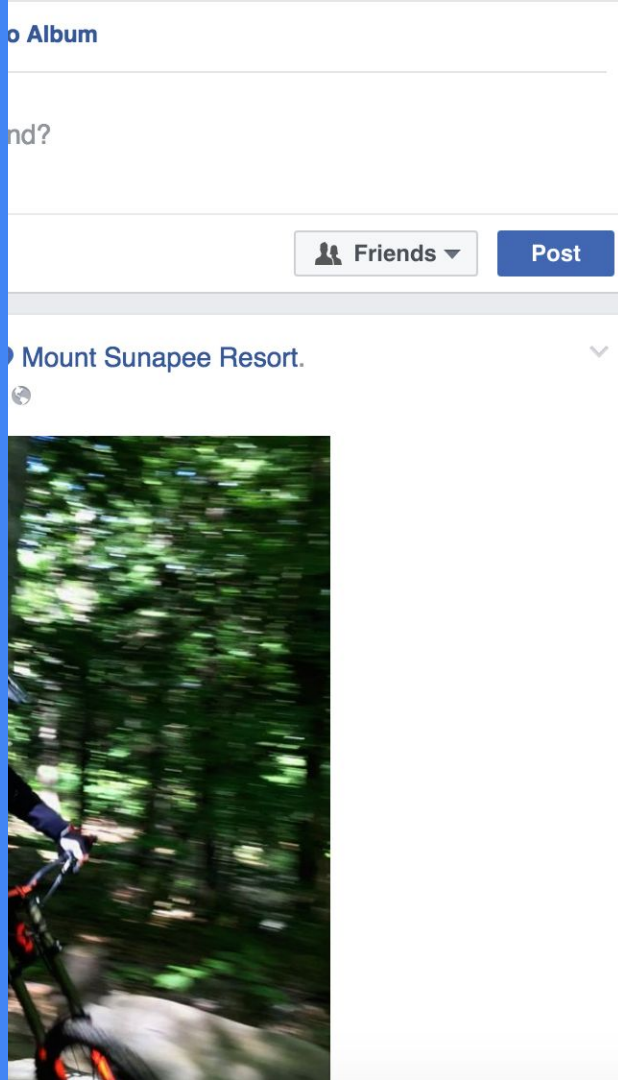
History of Big Data

- 2004 - Facebook was founded (yay)
- 2005 - Web 2.0
- 2007 - iPhone and AWS was released (yay)
- 2008 - 14.7 EB of data being generated
- 2014 - Mobile system surpass desktops



Web 2.0

- Front End
 - AngularJS, etc
- Dynamic Content
 - PHP
 - JSP
 - ASP.NET
 - Ruby
 - JavaScript
- Data Processing
 - Hadoop, etc
- Database
 - NoSQL
- Users Generate Content



 **Navin Bhatia's** birthday is today




TODAY'S GAMES



7:10pm

TRENDING



-  **iPhone 7 Plus:** YouTube Group Releases Video Claiming to Detail Upcoming Device
-  **Norah Jones:** Singer Releases Music Video for New Single 'Carry On' Off Upcoming Album
-  **Frank Ocean:** Singer's 'Boys Don't Cry' Album Reportedly to Be Released Friday on Apple Music

▼ See More

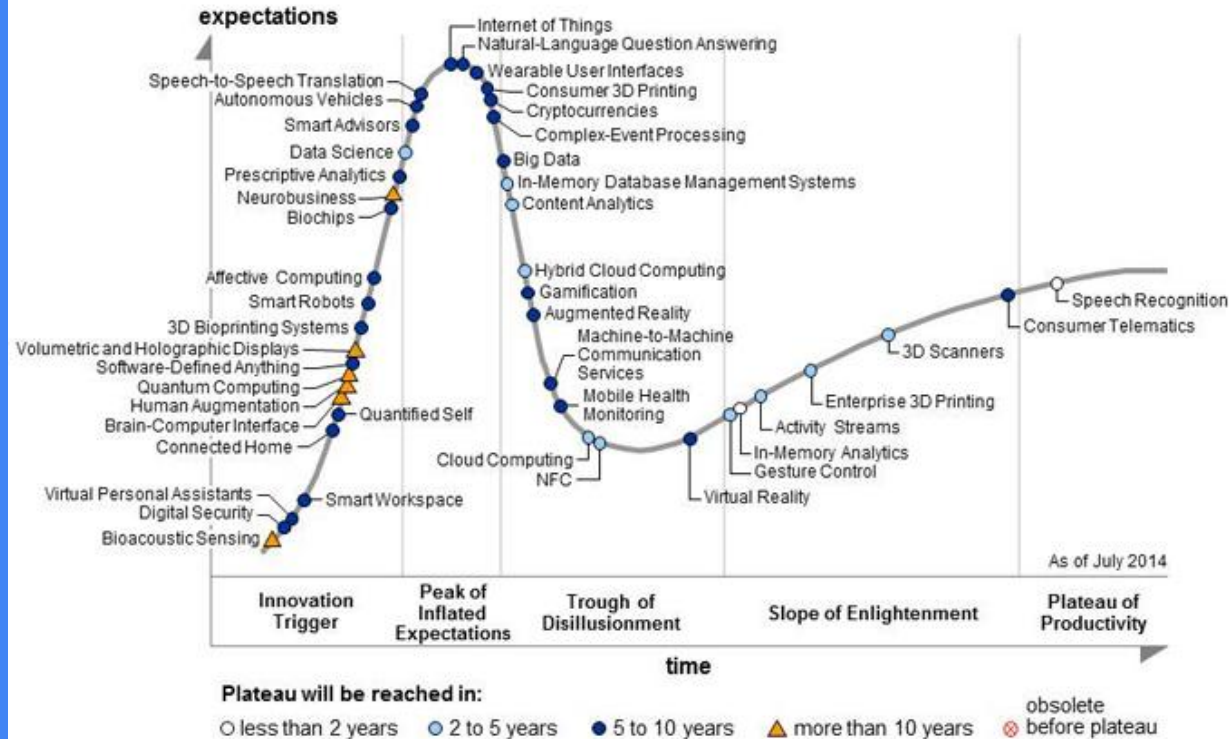
SPONSORED

Create Ad



Gartner Cycle of Hype 2014

- Every year Gartner publish a list buzzwords and their trends
- Divided into 5 categories
- Divided into 5 time ranges
- Big Data is doing OK



Gartner Cycle of Hype 2015

- Big Data is gone!?



History of Big Data

- 2014 - end of a buzzword
- 2015 - the wide deployment stage





More on Big Data

- History of Big Data
- **Frameworks and Tools**

Explosion of Frameworks and Tools



Frameworks and Tools

- Google leads the way

- Google File System
- MapReduce
- Bigtable
- Chubby
- Pregel
- Dremel
- Tenzing
- Spanner
- F1
- Borg
- Omega



Frameworks and Tools

- Community follows
 - Apache Hadoop HDFS
 - Apache Hadoop MapReduce
 - Apache HBase
 - Apache Zookeeper
 - Apache Pig
 - Apache Hive
 - Apache Drill
 - Apache Impala
 - Apache Giraph
 - Apache Mesos
 - Apache Spark





Key Projects

- Apache Zookeeper
- Apache Hadoop
- Apache Kafka
- Apache Mesos

Apache Zookeeper

- Make building distributed system easier
- Enabled a wide range of open-source projects
- Ground of truth for distributed system



Apache Hadoop

- Make distributed computation easier
- Expose simple API
 - Map
 - Reduce



Apache Kafka

- Make data transportation easier
- Make SOA model more reliable
- De facto data transportation framework



Apache Mesos

- Make distributed task scheduling easier
- Dramatically increase resource utilization




Frameworks and Tools

- Computation
 - Hadoop, Spark, Samza, Flink, Hive, Pig, Drill, etc
- Transportation
 - Kafka, Flume, Sqoop, Scribe, RabbitMQ, ZeroMQ, IronMQ, etc
- Storage
 - HBase, Cassandra, CouchDB, MongoDB, etc
- Coordination
 - Zookeeper, Consul, Etcd, Eureka, etc
- Scheduling
 - Mesos, Yarn, Oozie, etc

Explosion of Frameworks and Tools

- Don't Panic
- Building Big Data Platform become LEGO building





Introduction to Zookeeper

Because distributed system is a
ZOO



Agenda

- **Use Cases**
- What is Zookeeper
- Architecture
- Zookeeper Usage

An Example Computing Problem

Given a truck of fruits, calculate the quantity for each kind of fruit



Standalone Program

- Runs on a single machine
- Predictable
- Isolated
- Performance is limited by the machine




Distributed System/Program

- Runs on multiple machines
- Distribute the workload
- Hard to coordinate



Master Slave Model

- Master knows a list of tasks
- Master in charge of distributing tasks
- Slaves in charge of executing tasks
- Slaves need to report status back to Master



Master Slave Model Common Issues

- How to come up with a Master?
- What happens when master crash?
- What happens when worker crash?
- What if master and worker cannot communicate?



Master Slave Model Common Tasks

- Master Election
 - The process of deciding who is the master
- Crash Detection
 - The master need to detect when workers crash
- Group Membership
 - The master must learn who is available for tasks
- Metadata Management
 - All the nodes must be able to reliably store/retrieve status





Agenda

- Use Cases
- **What is Zookeeper**
- Architecture
- Zookeeper Usage

What is Zookeeper

- An open source distributed system that provides
 - Strong consistency, ordering, and durability guarantees
 - The ability to implement typical synchronization primitives
 - A simpler way of dealing with concurrency
- Inspired by Google Chubby
- Developed in Yahoo using Java





Agenda

- Use Cases
- What is Zookeeper
- **Architecture**
- Zookeeper Usage



Architecture

- **Coordinate Strategy**
- Concepts
- Internal


Ways to Coordinate

- Message passing
 - Processes exchange messages directly through a network
- Shared storage
 - Read or write to shared storage



Coordinating with Messages

- Communicate to each other by passing messages



Coordinating with Message Passing

- Network Delays
 - Messages might get delayed arbitrarily
- Speed of Processing
 - Some machines might process things faster
- Clock Difference
 - The time on the machines are different
- Cannot tell the difference between real crash or network issue



A Split-brain Example

- In a system, we have one master, one backup master, and many workers
- Part of the workers are having trouble to reach master
- Connect with backup master instead



Coordinate with Shared Storage

- More straightforward
- Still rely on network to transfer data
- Zookeeper use this model to implement coordination



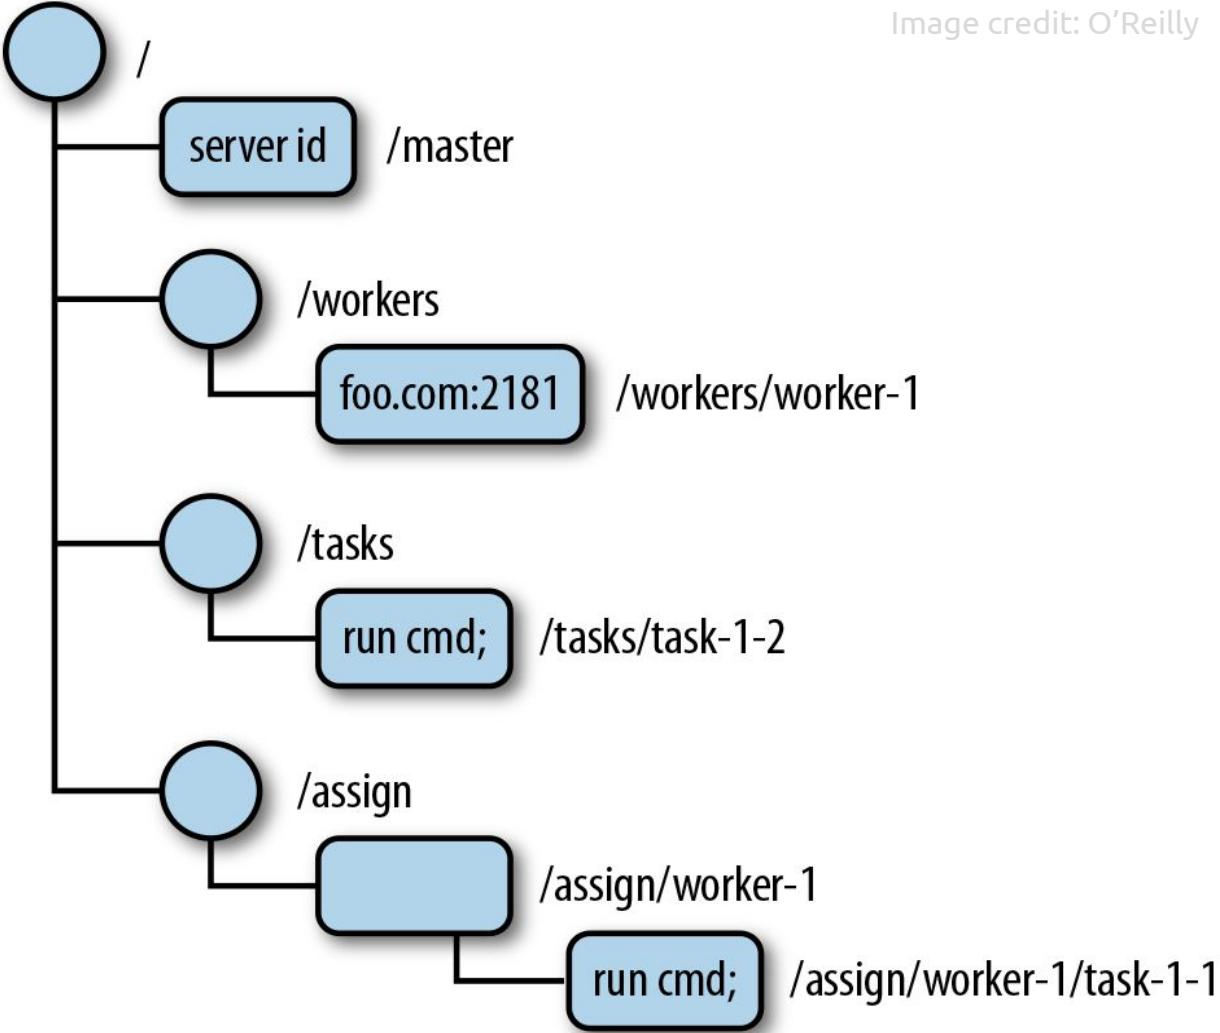


Architecture

- Coordinate Strategy
- **Concepts**
- Internal

Data Tree

- Organized in hierarchical structure
- Similar to file system



API

API Name	Usage
create /path data	Create a znode /path with data
delete /path	Delete znode /path
exists /path	Check if /path exists
setData /path data	Set znode /path to data
getData /path	Get the data in /path znode
getChildren /path	Return a list of children under /path



znode

- Basic unit for Data Tree
- Persistent znode
 - Data is persisted unless `delete` is called
- Ephemeral znode
 - Node is deleted if the client created it loses connection to zookeeper
- Sequential znode
 - Zookeeper will assign sequence number and append it to the path




Watcher

- Help client to know changes to znodes
- Avoid race conditions in polling
- Notifications are one-time operation



Watcher





Architecture

- Coordinate Strategy
- Concepts
- **Internal**

Client - Server Interaction

- Connection
 - TCP connection
 - Client only connects to server with newer/equal state
 - Configurable session timeout
- Read
 - Read can be read from any of the servers
- Write
 - Write requests will be forwarded to the leader of zookeeper cluster



zxid

- A 64-bit integer id
- Generated when there is a new update of znode

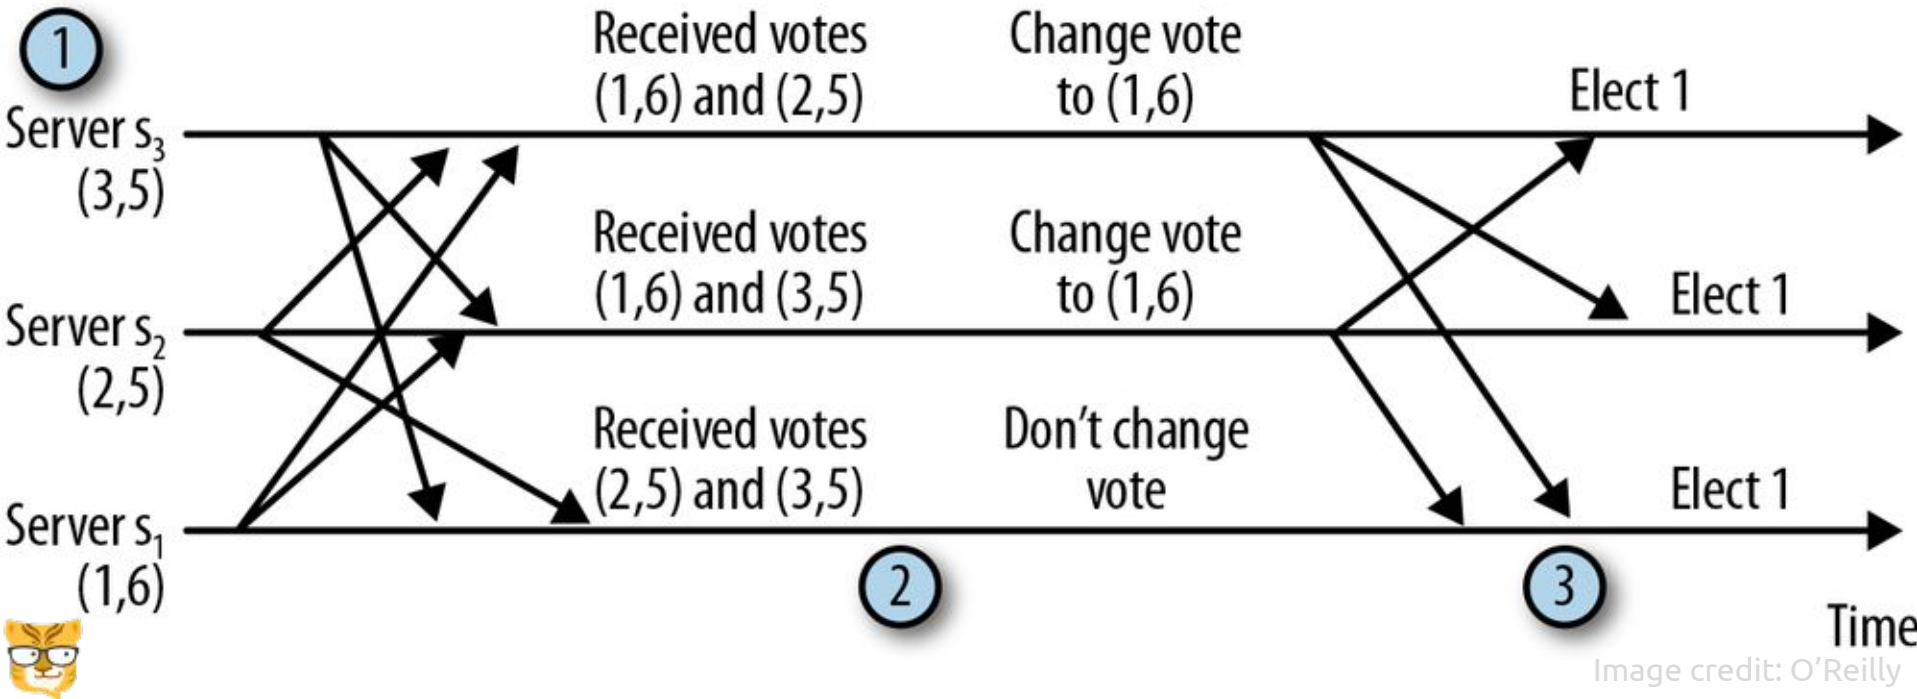


Leader Election within Zookeeper

- Zookeeper cluster AKA Ensemble
- Server has following mode
 - Leader
 - Follower
 - Observer
- Server has following state
 - LOOKING
 - LEADING
 - FOLLOWING



Leader Election within Zookeeper



State Replication within Zookeeper

- Follow a consensus protocol called Zab
- Leader
 - Receives Write requests
 - Convert requests into transaction
 - Leader send PROPOSAL message to all follower
 - Once receiving acknowledges from a quorum, leader sends COMMIT message



State Replication within Zookeeper

- Follower
 - Received PROPOSAL message from leader
 - Check if the leader is the correct leader
 - Check if the transaction is in correct order
 - Send acknowledgement back to leader
 - Received COMMIT message from leader
 - Apply change to the Data Tree





Agenda

- Use Cases
- What is Zookeeper
- Architecture
- **Zookeeper Usage**



Zookeeper Usage

- Master Election
- Crash Detection
- Group Membership
- Metadata Management

Master Election

- Have all the process go and create sequential ephemeral znode
- Znode with smallest sequence number is the leader
- Setup watcher for changes



Crash Detection

- Have slaves create ephemeral znodes
- Setup watcher on the znodes



Group Membership

- The members need to create ephemeral znodes under group node



Metadata Management

- Store metadata in persistent znode
- Use persistent znode as source of truth



Similar Systems

- Consul (<https://www.consul.io/>)
 - Use Raft for Consensus
- Etcd (<https://coreos.com/etcd/>)
 - Use Raft for Consensus
- Eureka (<https://github.com/Netflix/eureka>)



Introduction to Kafka

A High-throughput Distributed
Messaging System



Agenda


- **Use Cases**
- What is Kafka
- Architecture
- Kafka Usage

Microservice Architecture

A monolithic application puts all its functionality into a single process...




... and scales by replicating the monolith on multiple servers



A microservices architecture puts each element of functionality into a separate service...




... and scales by distributing these services across servers, replicating as needed.



Microservice Architecture

- Organize around business capabilities
- Decentralized data management
- Better deployment infrastructure



microservices - application databases




How Does Microservices Communicate

- Sync
 - RESTful API
 - RPC frameworks
- Async
 - Drop a message and back to work



Message Queue



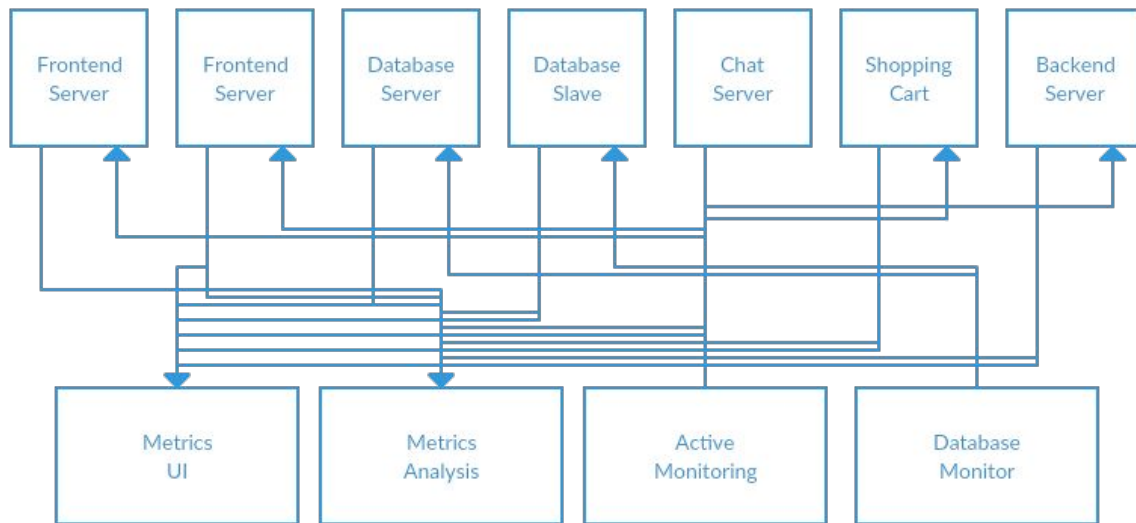
Message Queue Benefits

- Service Decoupling
- Increase Scalability
- Data Redundancy
- Deal with Peak Traffic
- Buffer Cushion for Failed Components



Message Queue Issues


- Spaghetti of queues
- Skip Messages
- Brokers need to record client positions



Application instances -
generating messages



Consumer service
instance pool -
processing messages





Agenda

- Use Cases
- **What is Kafka**
- Architecture
- Kafka Usage

What is Kafka

- An open source distributed messaging system
 - Fast - hundreds MB/s from thousands of clients
 - Scalable - easily scale up and down without downtime
 - Durable - Messages are persisted on disk to prevent data loss
- Developed in LinkedIn using Scala





Agenda

- Use Cases
- What is Kafka
- **Architecture**
- Kafka Usage



Architecture

- **Design Strategy**
- Concepts
- Internal

Pull vs Push

- Push model
 - High throughput
 - Complex server logic
- Pull model
 - Simple server logic
 - Reply feature



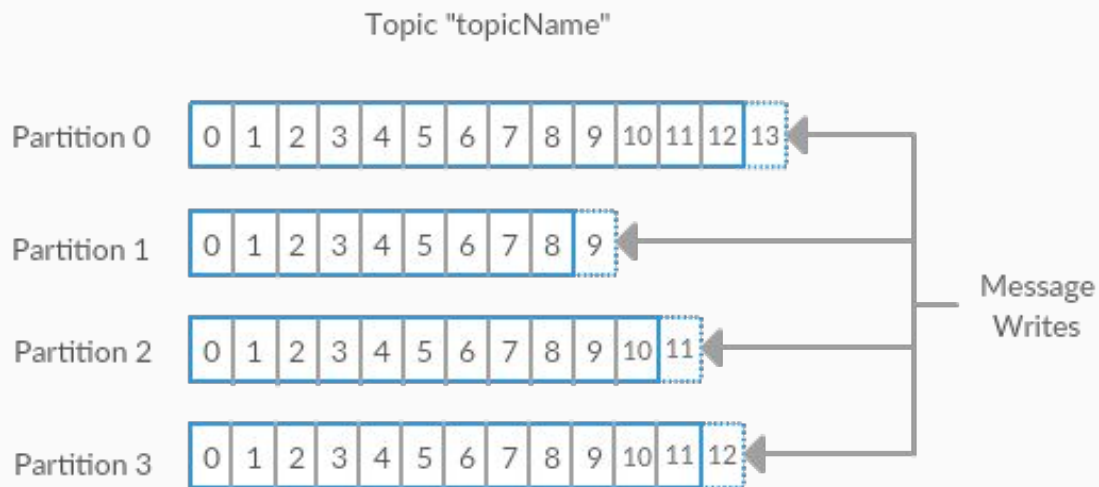


Architecture

- Design Strategy
- **Concepts**
- Internal

Topic and Partition

- Messages are organized logically into topics
- Physically divided into Partitions



Offset

- A incremental sequence number
- The position of a message in a partition



API

API Name	Usage
<code>publish topic data</code>	Publish data onto a topic
<code>consume topic offset</code>	Consume from a topic



Producer

- Producer in charge of sending messages to Kafka broker



Consumer

- Consumer pull messages from Kafka broker
- Can use offset to target a specific message on a partition
 - By default point to the latest offset
 - Can set the offset to an older value to read old data
- Consumer maintain message state
- No ACK



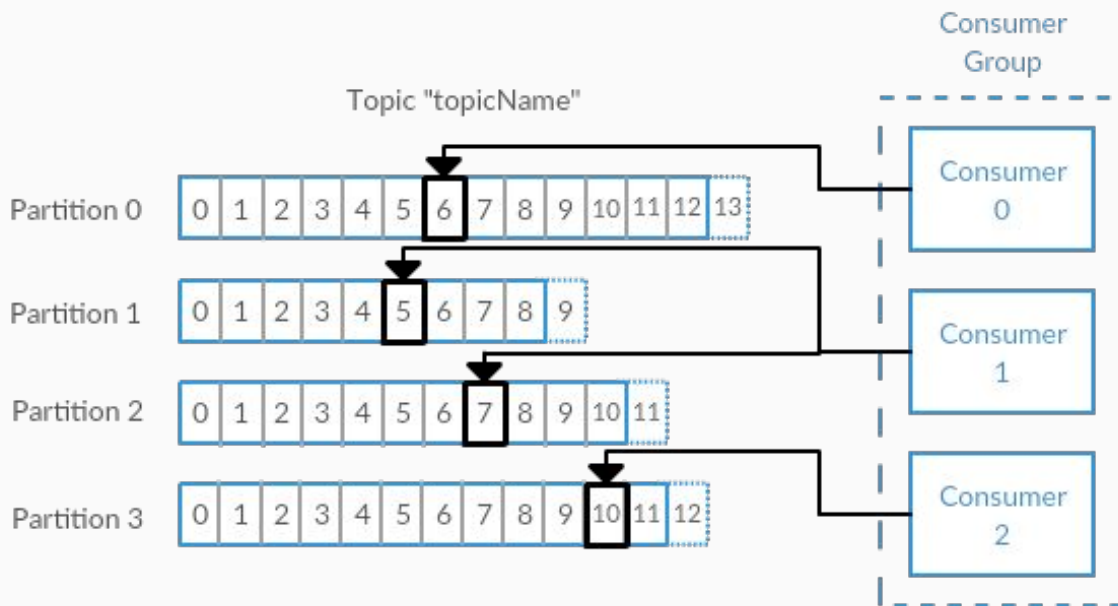
Consumer Group

- A group of consumers
- Mapped to partitions
 - Messages in certain partition can only be consumed by corresponding consumer



Consumer Group

- A group of consumers
- Mapped to partitions
 - Messages in certain partition can only be consumed by corresponding consumer





Architecture

- Design Strategy
- Concepts
- **Internal**

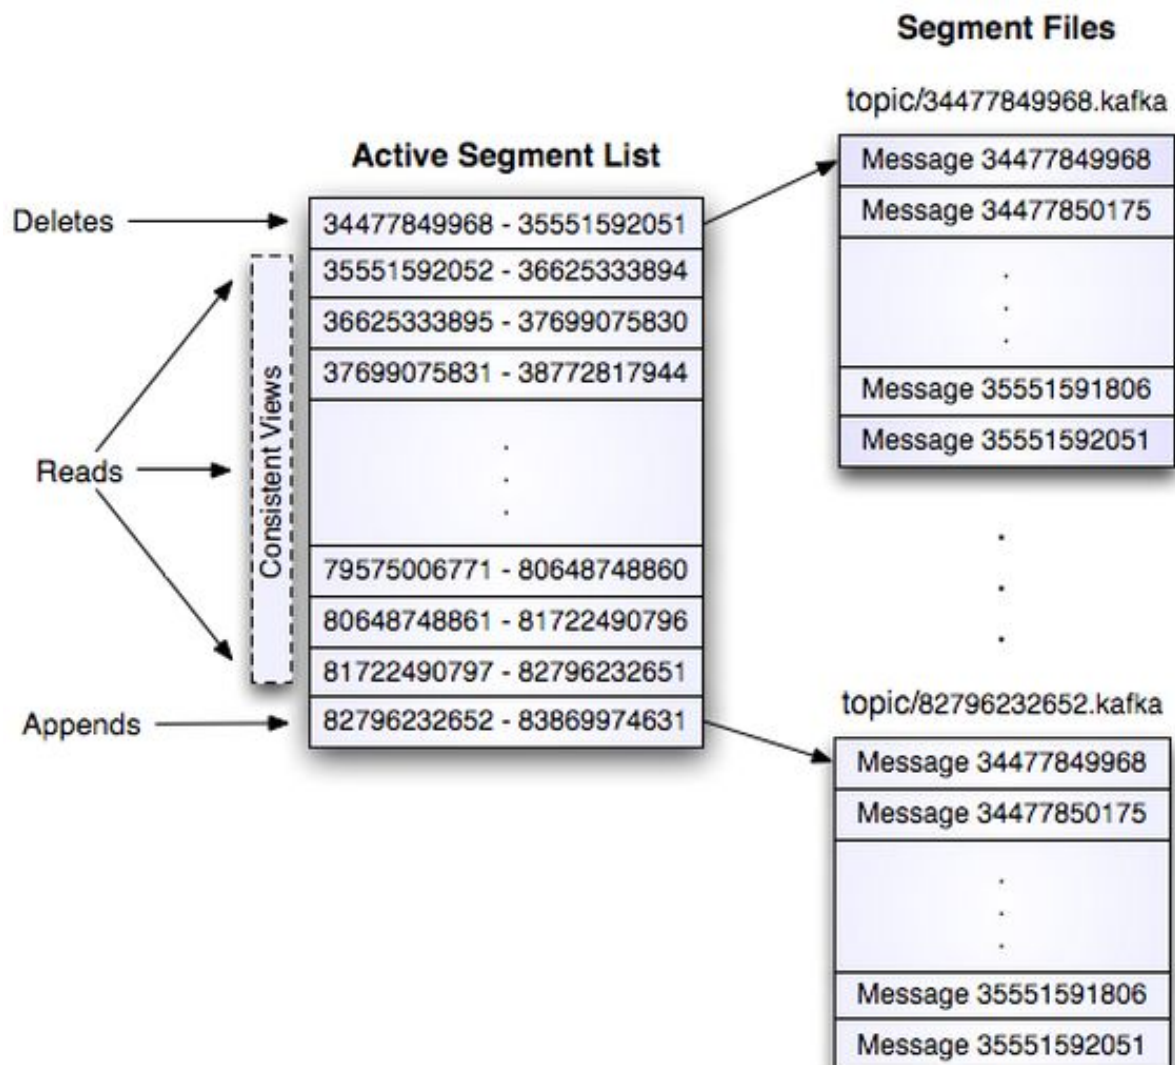
Log File Format

- One partition is a physical folder
- One message:
 - offset
 - Message length
 - Magic value, 1 byte
 - CRC value, 4 bytes
 - Payload, N bytes




Log File Format

- Active Segment List maintains mapping to segment files
- Log entries in segment files



IO Optimization

- Append-only writing
- Reads do not block writes
- Super fast writing speed because one partition is one log file



IO Optimization - zerocopy

- OS reads data from disk into pagecache in kernel space
 - Application reads data from kernel space into user space
 - Application writes data back to kernel space into socket buffer
 - OS copies data from socket buffer to NIC buffer
-
- zerocopy copies data into pagecache only once and reuse



Data Replication inside Kafka

- Producer write through Partition Leader
- Partition Leader write the message into local disk
- Partition Follower pull from Partition Leader
- Once Leader received ACK from all the



ISR (in-sync Replica)

- A nice balance between sync and async copy
- Configurable lag behind
 - If one replica is too slow, will be removed from ISR
- Follower can batch read from Leader





Agenda

- Use Cases
- What is Kafka
- Architecture
- **Kafka Usage**



Kafka Usage

- Log Aggregation
- Real-time Data Analysis

Log Aggregation

- Forward all the logs into specific Kafka topic
- Setup Consumer/Consumer Group on the topic and forward to indexing service
- Query on the data through Kibana, etc



Read-time Analysis

- Forward interested data into Kafka topic
- Integrate with Spark/Samza for steaming processing



Similar Systems

- RabbitMQ (<https://www.rabbitmq.com/>)
 - Erlang, based on AMQP
- ZeroMQ (<http://zeromq.org/>)
 - Middleware less
- Redis (<http://redis.io/>)
 - C/C++, can be used as lightweight queue
- ActiveMQ (<http://activemq.apache.org/>)
 - Java, based on AMQP
- SQS (<https://aws.amazon.com/sqs/>)
 - An AWS service, the visibility timeout feature is awesome





Introduction to Cassandra

Manage massive amounts of
data, fast, without losing sleep



Agenda


- **Use Cases**
- What is Cassandra
- Architecture
- Cassandra Usage

An Example Storage Problem

- We are creating a system for people to store:
 - First name
 - Last name
 - Phone number
- Initially we can just go with one simple table


An Example Storage Problem

- Your application is getting popular, you noticed that people are getting slow read
- Read/Write separation



An Example Storage Problem

- Your app is so popular you are hitting the storage limit of database
- Shard/Partition data based on some rules





Agenda

- Use Cases
- **What is Cassandra**
- Architecture
- Cassandra Usage

What is Cassandra

- An open source distributed storage system that provides
 - High availability
 - No single point of failure
- Inspired by Amazon DynamoDB
- Developed in Facebook using Java





Agenda

- Use Cases
- What is Cassandra
- **Architecture**
- Cassandra Usage




Architecture

- **Design Strategy**
- Concepts
- Internal

Revisit Master Slave Model

- Single point of failure
- Even backup master might fail
- Capacity depends on master



Gossip Protocol



- Runs every second
- Choose a random node to gossip with
- Use accrual failure detection to determine if a node is down
 - For example, increase the convict threshold on cloud services

Tunable Consistency

- Consistency means whether read always return the most recently written value
- In other systems, the consistency level is defined by the protocol



CAP Theorem






Architecture

- Design Strategy
- **Concepts**
- Internal

Ring

- In a Cassandra cluster, data is assigned to nodes as if they form a ring of tokens



Partitioner

- A hashing algorithm to determine how data is distributed across the cluster
- By default murmur3 hashing algorithm is being used

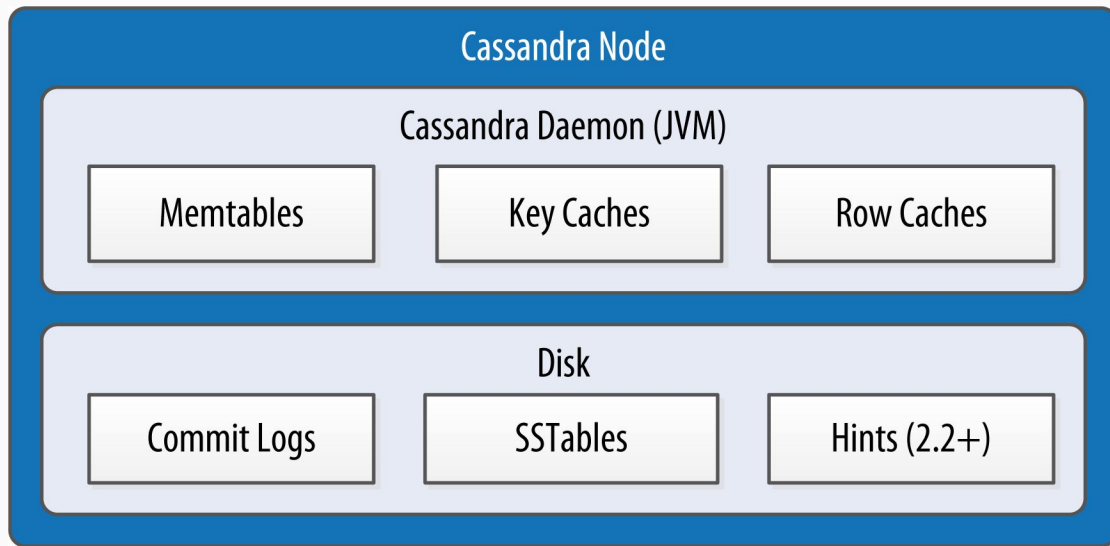


Architecture

- Design Strategy
- Concepts
- **Internal**

Internal Data Structure

- Commit log
- memtable
- SSTable



Commit log

- Crash-recovery mechanism
- All write operation is immediately written to commit log
- Will not count if no commit log is written

memtable

- Value will be added to memtable after commit log
- In-memory store to speed up operations

SSTables

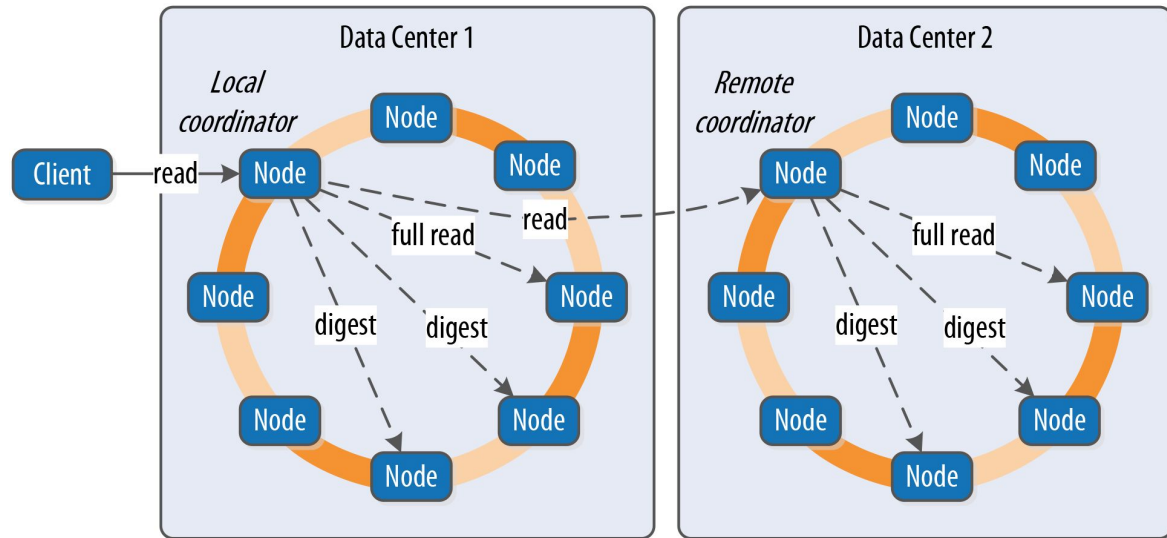
- Content of memtable gets written to SSTable after memtable is full
- Immutable cannot be changed
- Changes are appended
- Sequential write to disk

Compaction

- Merge of SSTables
- New merged data is sorted as well
- Reduce number of seeks


Read Operations inside Cassandra

- Client can contact any node to read




Write Operations inside Cassandra

- Client can contact any node to Write



Data Replication Inside Cassandra

- A node can serve as replica for other ranges
- If nodes goes down replica will serve the request





Agenda

- Use Cases
- What is Cassandra
- Architecture
- **Cassandra Usage**



Cassandra Usage

- General Data Storage
- Time-series Data Storage
- TTL Data Storage

General Data Storage

- You can use Cassandra cluster as your primary data persistence layer
- Easy to operate



Time-series Data Storage

- Data is sorted and written sequentially to disk
- Perfect for retrieving data and filter range
- Fast access due to small disk seeks



TTL Data Storage

- Some data can be discarded after some time
- With Cassandra TTL on data, this feature is easy to implement



Similar Systems

- HBase (<https://hbase.apache.org/>)
 - Java, inspired by Google BigTable
- MongoDB (<https://www.mongodb.com/>)
 - C++, C, JavaScript, document based database
- CouchDB (<http://couchdb.apache.org/>)
 - Erlang, document based database





Q&A