



Big Data Engineer Bootcamp

Part 2



Agenda

- Introduction to Docker
- Introduction to Redis
- Introduction to Node.js
- Introduction to Spark
- Introduction to Mesos
- Q&A



Introduction to Docker

Build, Ship, and Run Any App,
Anywhere



Agenda

- **Use Cases**
- What is Docker
- Architecture
- Docker Usage

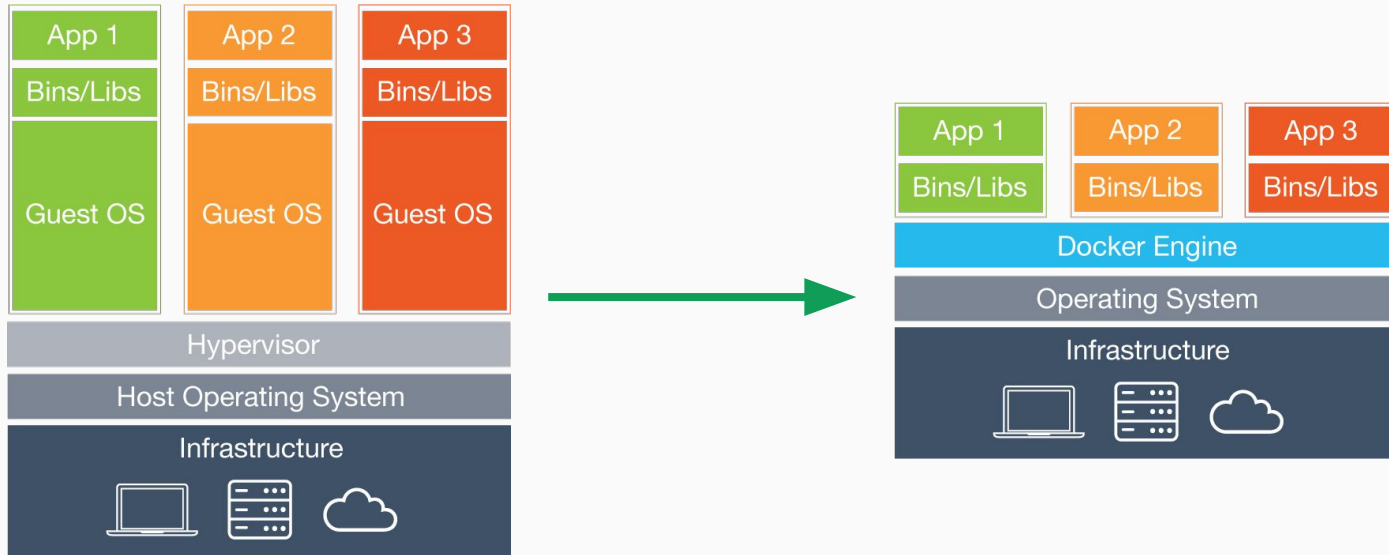
Consistent Deployment Model

- In the industry, a software will go through multiple stages
 - Development machine
 - QA - testing
 - Staging - integrate with other teams
 - Production - handles real traffic

	Development Machine	QA Environment	Staging Environment	Production Environment
data-producer.py	?	?	?	?



Increase Resource Utilization



Docker Benefits

- Faster developer onboarding
- No vendor lock-in
- Eliminate environment inconsistencies
- Ship applications faster
- Scale quickly
- Easily remediate issues





Agenda

- Use Cases
- **What is Docker**
- Architecture
- Docker Usage

What is Docker?

- A tool to package and deploy applications inside containers
 - Containers are isolated environments
- Developed by Solomon Hykes in Dotcloud
- Open-sourced in March 2013, written in Go
- Grown into a platform
 - Docker Compose, Docker Swarm
 - Docker Image Hosting
 - Container Hosting





Agenda

- Use Cases
- What is Docker
- **Architecture**
- Docker Usage
- Hands on

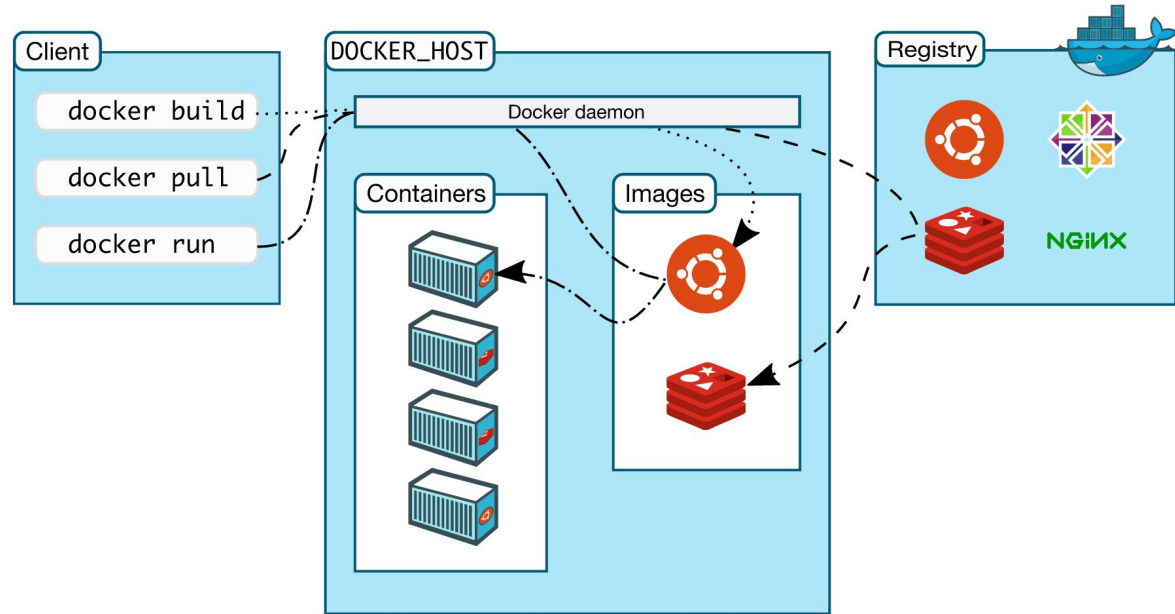


Architecture

- **Concepts**
- Internal

Client Server Architecture

- Docker Client
- Docker Daemon
- Docker Registry



Docker Client

- A small client to communicate with Docker Daemon
 - User interact with Docker Client to perform tasks
- Can connect to any remote Docker Daemon



Docker Daemon

- A background daemon running on host servers
 - a small server takes requests from Docker Client and dispatch/route to corresponding handler
 - an engine takes the requests and manipulates containers such as creating container, etc



Docker Registry

- A warehouse for container images
 - Similar to github
- Talks to Docker Daemon to handle image related requests from Docker Client
- You can use public Docker Registry (which is [Dockerhub](https://hub.docker.com/)) or private Docker Registry





Architecture

- Concepts
- **Internal**

Docker Image

- The basis of containers
- Contains a number of layers
- Use union filesystem to combine layers to form a single read-only file system
- Docker Images are built from
 - run a base container, run commands within the container and commit, push the changes like git
 - base images (such as ubuntu) and set of instructions
- Instructions are stored in a file called Dockerfile




Docker Image

- In Docker, each layer is described by:

- Meta data of the layer, in JSON format
- Image Filesystem changeset
- Image ID, such as:

a9561eb1b190625c9adb5a9513e72c4dedafc1cb2d4c5236c9a6957ec7dfd5a9



Docker Image

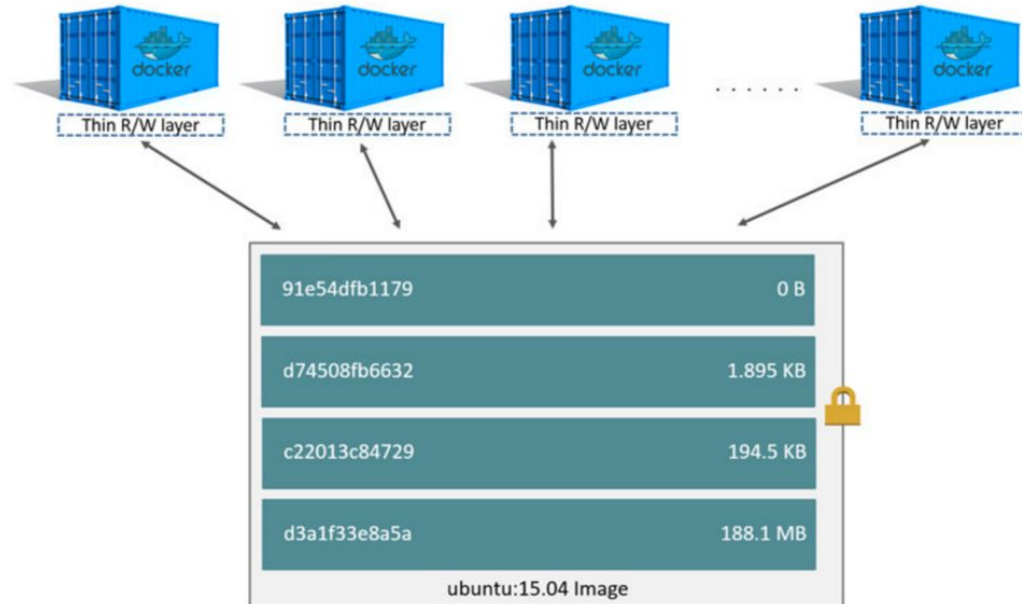
- Look inside `/var/lib/docker/aufs/layers`
- Look inside `/var/lib/docker/aufs/diff`

```
f4c34dfa6c6b8e364fce3a8b15da94a4f33e078cbad3927140f7dc78d6a6ad30
f693368ae907e53a4ae5c2d875abe84af5c1bf10eaf6529e2718f5b0e73b165f
fa3839de4485f732f5394331e1d2e872c4bd34e11497c35f953c4b04832d0e76
fa3ea1829dfe1e2a86b0e0b0bc7899e8d7adfcf4aeacbd26bf7b0cf091bcdea7
fcbfeb515f08fadfdef9dc6d6b7f738c7196ccc4b06a875fcf4da06903c42b84
fcd053878fdb70ae2410cb686c8e2944ece7602d9ee985d5a6850cf00ab784b3
fd28c2fe6e450d39c6a74ee622de5e5f7017c5e4ee8025140e9c0570c2f2d73f
fdc46298686a74d04f09c708ce5f7007c0504dbbdb506dfb9bf8c0c4a7e3e5f8
fe21d21ea2f95678074b48e017a81cab951ff3a51b24feac93efabac2468056d
ff40d161449915fe26ce7d099f817c1185291b48814fad6c476e58433f2fe275
ff6e1984778b2f3e1ed13737499eda4f04cf347df6b36db42d9698764d072c67
root@default:/mnt/sda1/var/lib/docker/aufs/layers#
```

```
fd28c2fe6e450d39c6a74ee622de5e5f7017c5e4ee8025140e9c0570c2f2d73f
fdc46298686a74d04f09c708ce5f7007c0504dbbdb506dfb9bf8c0c4a7e3e5f8
fe21d21ea2f95678074b48e017a81cab951ff3a51b24feac93efabac2468056d
ff40d161449915fe26ce7d099f817c1185291b48814fad6c476e58433f2fe275
ff6e1984778b2f3e1ed13737499eda4f04cf347df6b36db42d9698764d072c67
root@default:/mnt/sda1/var/lib/docker/aufs/diff# cd ff6e1984778b2f3e1e
d9698764d072c67
root@default:/mnt/sda1/var/lib/docker/aufs/diff/ff6e1984778b2f3e1e
8764d072c67# ls
opt
root@default:/mnt/sda1/var/lib/docker/aufs/diff/ff6e1984778b2f3e1e
8764d072c67#
```



Docker Container



graphdb

- Manage local Docker Image and their relationships
- Internally use SQLite




Dockerfile Example

```
FROM mhart/alpine-node:base
ADD . .
EXPOSE 3000
CMD ["node", "index.js"]
```



namespace

- A linux kernel functionality that can
 - perform resource limiting
 - resource prioritization
 - resource accounting
- Docker uses cgroup to achieve resource limiting



cgroup

- A linux kernel functionality that can
 - perform resource limiting
 - resource prioritization
 - resource accounting
- Docker uses cgroup to achieve resource limiting



Docker Driver

- Customize the execution environment of Docker Containers
- **graph driver** - related to storage
 - aufs, devicemapper, btrfs, zfs, and overlay
- **network driver** - related to network
 - bridge, ip, port
- **exec driver** - related to the execution of container
 - LXC, runC



runC

- An abstraction layer between Docker Driver and Linux kernel
- It interfaces with Linux kernel functionalities such as:
 - namespaces
 - cgroups
 - capabilities
 - file system access controls





Agenda

- Use Cases
- What is Docker
- Architecture
- **Docker Usage**



Docker Usage

- Speed Up Workflow
- Handle Spike Traffics

Speed Up Workflow

- Push from dev machine to production in seconds
- Deliver functionalities to end user multiple times a day
 - CI/CD + Docker

	Development Machine	QA Environment	Staging Environment	Production Environment
data-producer.py	!	!	!	!



Handle Spike Traffics

- Docker is perfect for stateless tasks/jobs
- Starts up within second



Further Reading

- Docker Cheatsheet: <https://github.com/wsargent/docker-cheat-sheet>
- namespace: https://en.wikipedia.org/wiki/Linux_namespaces
- cgroup: <https://en.wikipedia.org/wiki/Cgroups>





Introduction to Redis

Swiss Knife Data Structure

What is Redis

- Open source in-memory data structure store
 - Database
 - Cache
 - Message Queue
- Developed by Salvatore Sanfilippo
- Implemented in C, high performance
- Super clean API + data structure



redis



Supported Data Structure

- Strings
- Lists
- Sets
- Sorted Sets
- Hashes (think of it as map)
- Bitmap
- Hyperloglog



Supported API

- Create/Read/Update/Delete
 - SET key value
 - GET key
 - LPOP key
- TTL
 - SETEX key value ttl
- And many more (<http://redis.io/commands>)



Use Cases

- Use Redis as LRU cache
- Use Redis as non-critical message queue
- Basically think of Redis as your Leetcode as a server ;)



Further Reading

- Redis official documentation: <http://redis.io/documentation>
- Little Redis Book: <http://openmymind.net/redis.pdf>
- Salvatore Sanfilippo's blog: <http://antirez.com/latest/0> ****





Introduction to Node.js

Swiss Knife Data Structure

What is Node.js

- Open source platform for service-side web applications
 - Event-driven architecture
 - Async IO
- Developed by Ryan Dahl
- Instant popularity for stack cohesion
 - Emerge of full-stack engineer



Sample Node.js Web Application

- Functions are first level citizen in javascript
- Callback functions are heavily used in Node.js

```
main.js *  
  
// Load the http module to create an http server.  
var http = require('http');  
  
// Configure our HTTP server to respond requests.  
var server = http.createServer(function (request, response) {  
  response.writeHead(200, {"Content-Type": "text/plain"});  
  response.end("Bittiger\n");  
});  
  
// Listen on port 8000, IP defaults to 127.0.0.1  
server.listen(8000);  
  
// Put a friendly message on the terminal  
console.log("Server running at http://127.0.0.1:8000/");
```



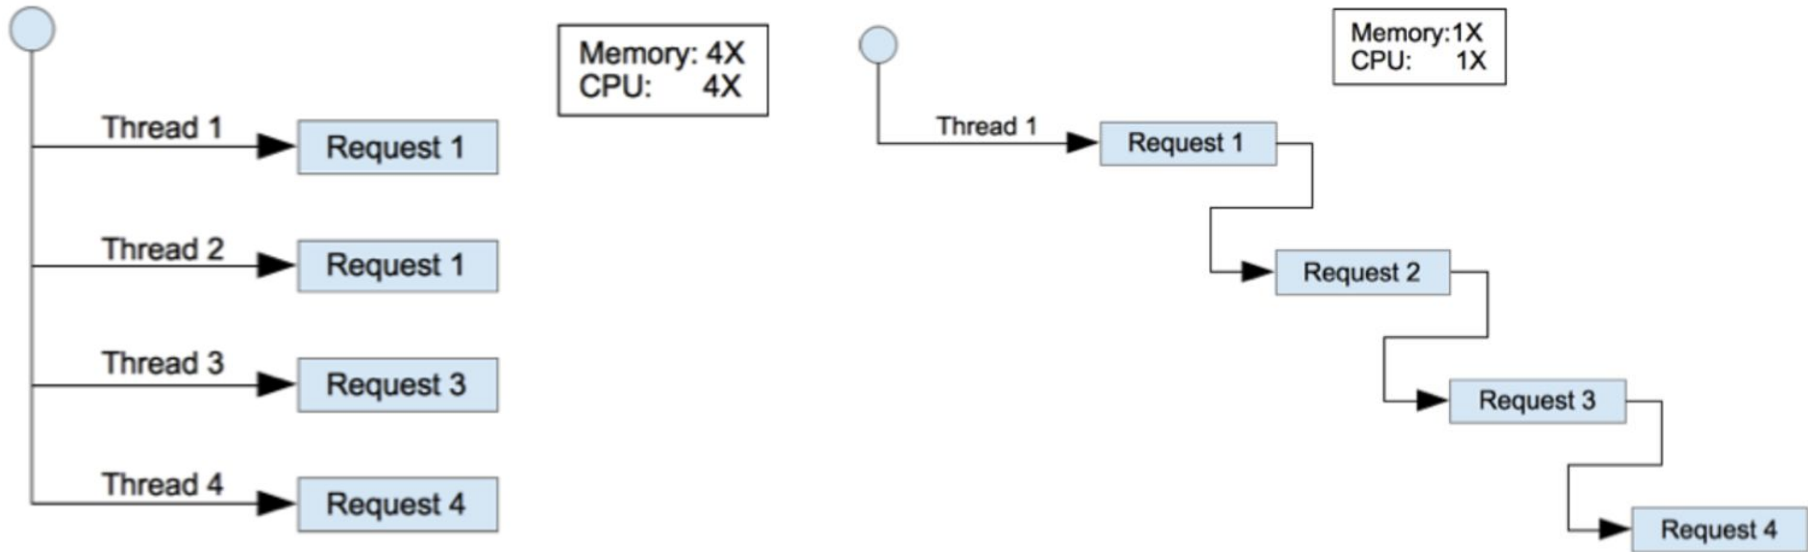
Comparison between J2EE and Node.js

- In web applications, majority of the work is IO

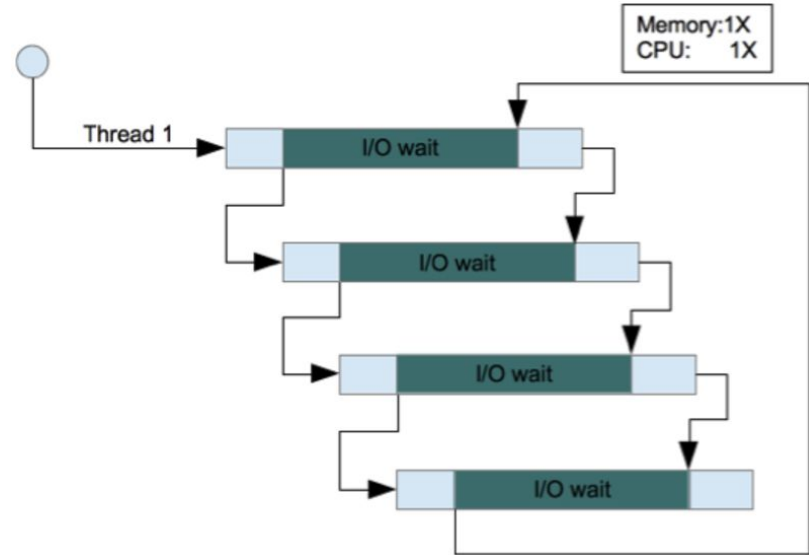

	Concurrency Model	I/O Model
J2EE	Multiple thread	Synchronous IO
Node.js	Single thread	Asynchronous IO



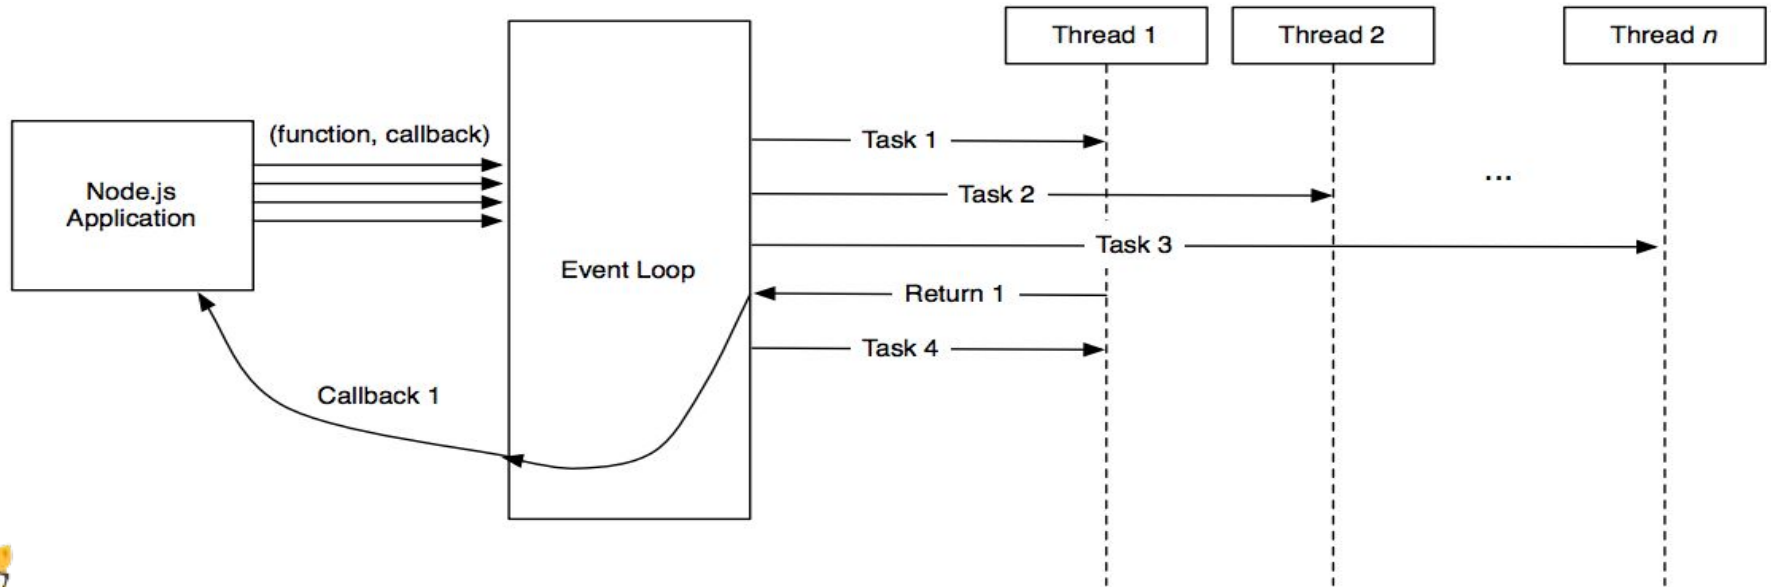
Multi-thread vs Single-thread



Synchronos IO vs Asynchronos IO



Eventloop



Further Reading

- Node.js Documentation: <https://nodejs.org/api/>
- Node.js Beginners Guide: <http://nodeguide.com/beginner.html>
- Deep dive into Node.js Architecture:
<http://www.journaldev.com/7462/node-js-architecture-single-threaded-event-loop>





Introduction to Spark

Lightning-fast Cluster
Computing



Agenda

- **Use Cases**
- What is Spark
- Architecture
- Spark Usage
- Hands on

Example Compute Problem

- Give a collection of fruits
- Count the quantity of each type of fruit



Example Compute Problem

- Naive Approach
 - For every type of fruit, find all of the items
 - Count the number of items
- Divide the Conquer
 - Split the fruits into small chunks
 - Count on the small chunks
 - Aggregate the results



Distributed Computing

- Divide and Conquer fits perfectly with Distributed Model
- Hard to implement
 - Network issue
 - Slave unavailable
 - No clean set of APIs



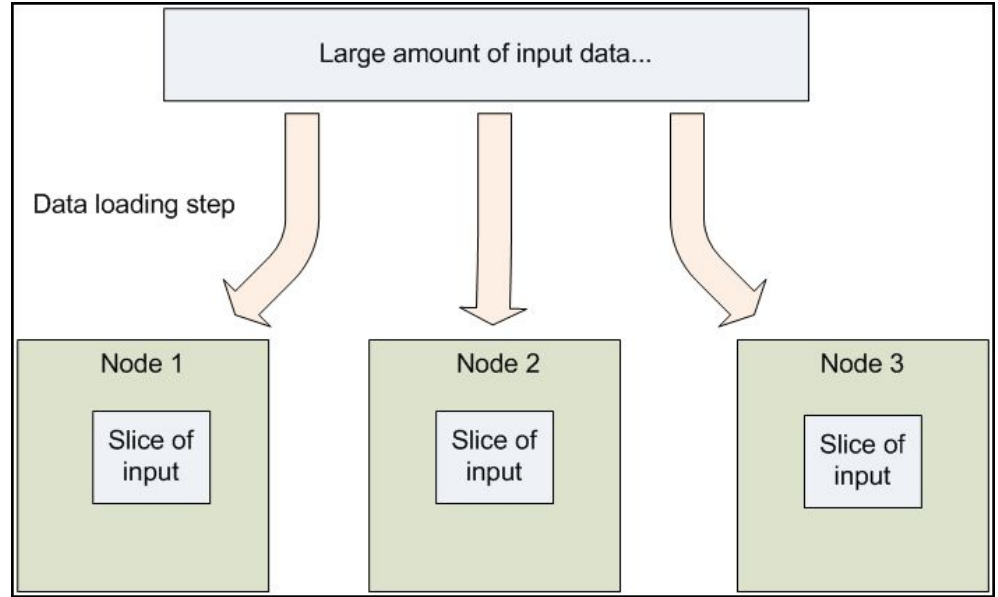
Apache Hadoop

- Open source distributed computing framework
 - Simply programming model
 - High resiliency to hardware failure/network failure
- Developed by Doug Cutting at Yahoo!
- Based on Google GFS and MapReduce paper
- Ground-breaking project from Apache Software Foundation



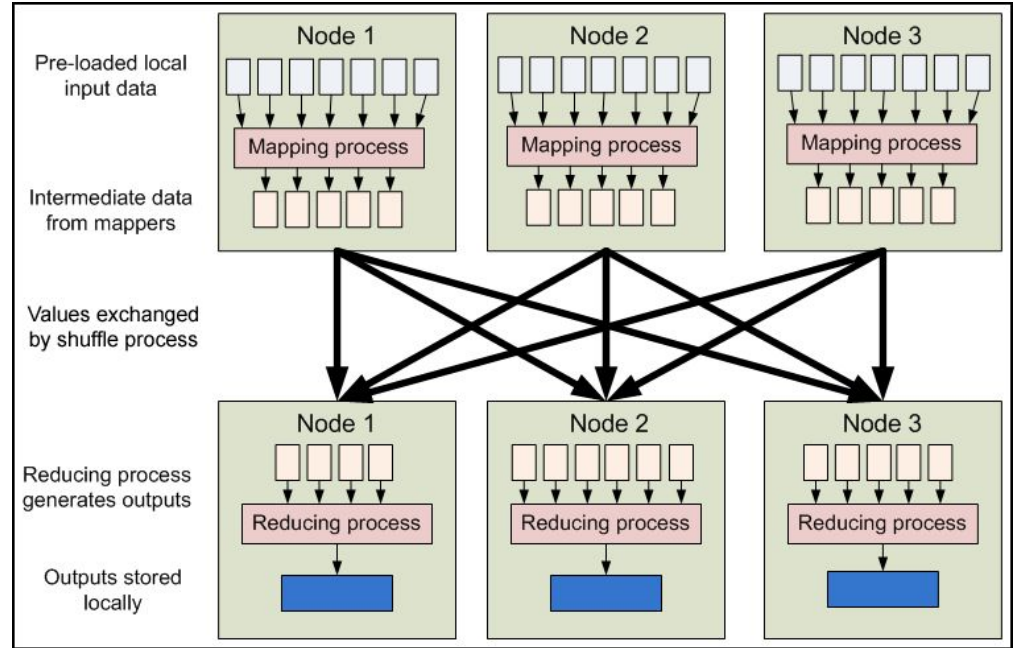
Apache Hadoop HDFS

- Handles data storage aspect of big data
- Master-slave architecture
- Data split into different slaves



Apache Hadoop MapReduce

- Slave nodes run computation on the sub-data set -> **MAP**
- Intermediate data are saved temporarily and then sorted
- Slave nodes run aggregation of sub-data set result -> **REDUCE**



Problem with Apache Hadoop

- Massive disk IO
 - Mapper write intermediate data into disk
 - Transferred through network for shuffling
 - Reducer load intermediate data from disk





Agenda

- Use Cases
- **What is Spark**
- Architecture
- Spark Usage
- Hands on

What is Spark

- Open source cluster computing framework
 - Respond to limitations of Apache Hadoop
 - Computation optimization
 - In memory computing
- Developed at UC Berkeley by Matei Zaharia

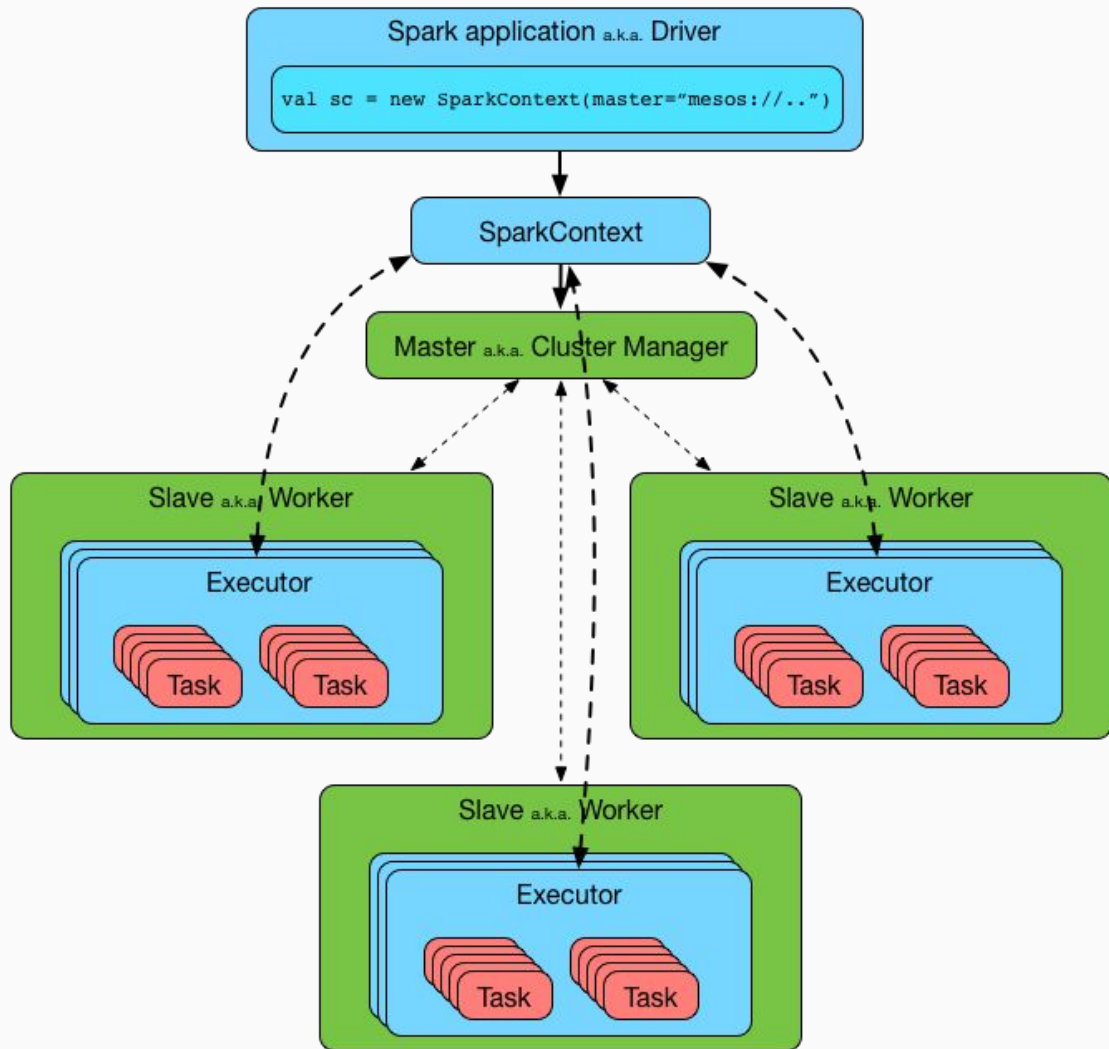




Agenda

- Use Cases
- What is Spark
- **Architecture**
- Spark Usage
- Hands on

Master Slave Mode

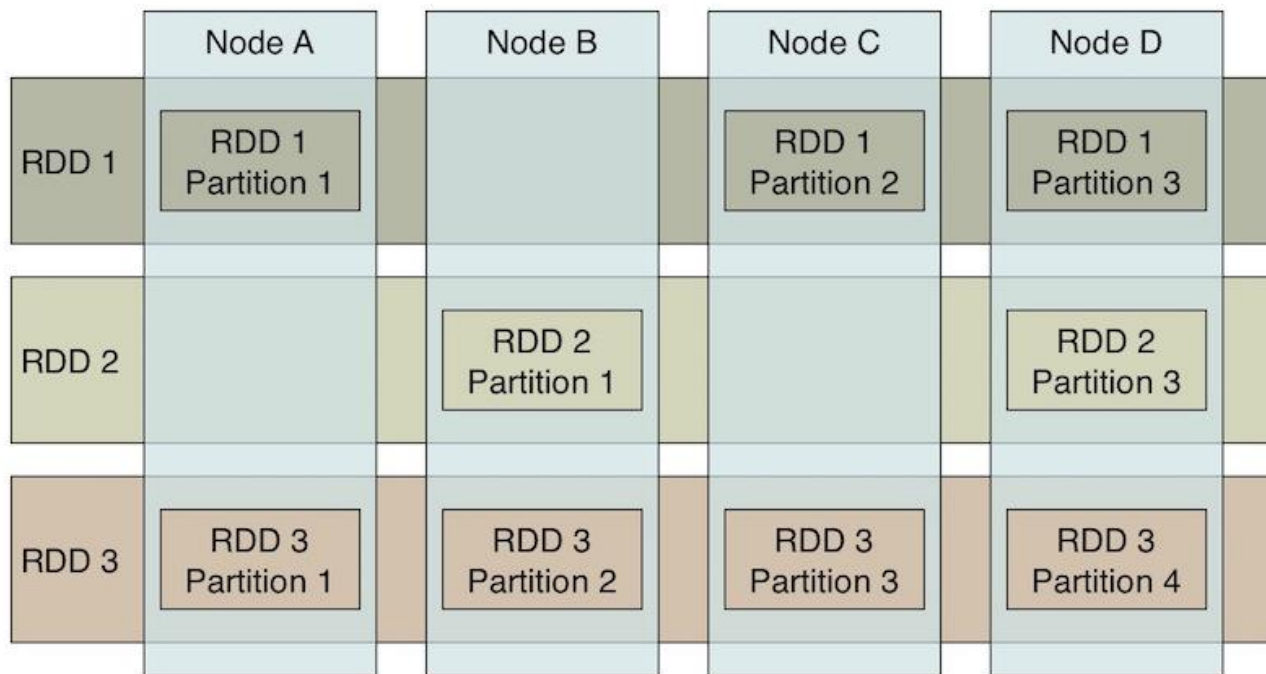


Resilient Distributed Datasets - RDD

- How Spark represents data
- RDD for one data set spread across the Spark cluster



Resilient Distributed Datasets - RDD



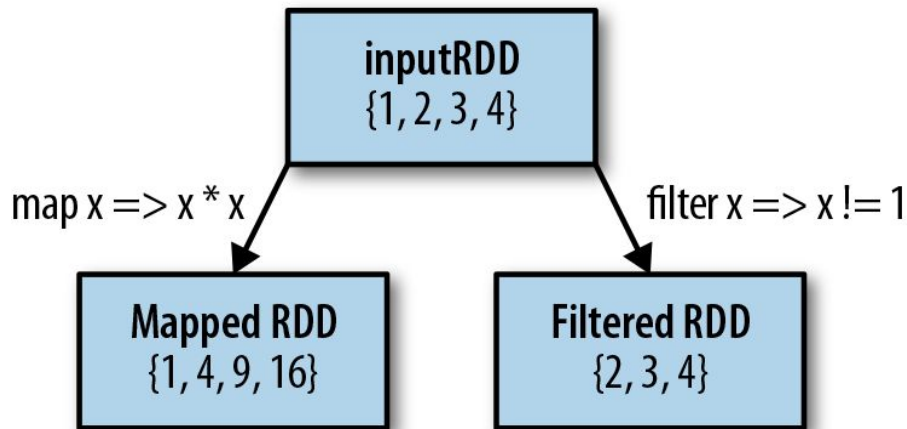
Resilient Distributed Datasets - RDD

- RDDs are immutable and readonly
- Can only be built by
 - Load data from raw storage
 - Transform from other RDD



RDD Transformation

- map
- filter
- flatMap
- mapPartitions
- sample
- union
- intersection
-



RDD Representation

- Each RDD has the following information
 - A set of partitions
 - A set of dependencies on parent RDD
 - A function for computing from its parents
 - Metadata about data placement



Lazy Evaluation

- Transformation won't actually perform calculation
- Aggregate compute steps for optimization
- Actual computation happens at action step



RDD Action

- collect
- count
- countByValue
- reduce
- top
- ...



RDD Persistence

- You can instruct Spark to cache certain RDD
- Configurable
 - Memory
 - Disk
 - Mix of both



Wordcount Example

- `text = sc.textFile('shakespeare.txt')`
- `counts = text.flatMap(lambda line: line.split(" ")).map(lambda word: (word, 1)).reduceByKey(lambda a, b: a + b)`
- `counts.collect()`





Agenda

- Use Cases
- What is Spark
- Architecture
- **Spark Usage**
- Hands on



Spark Usage

- Machine Learning Algorithms
- Stream Processing

Machine Learning Algorithms

- Many Machine Learning Algorithms runs iteratively on data sets
- Perfect to run on Spark
 - Persist/Cache RDD
 - DAG computation step optimization



Stream Processing

- The Spark Stream library
 - Provides API to run computation over data stream
 - Integration with common data sources
 - Kafka
 - Flume
 - HDFS
 - S3





Introduction to Mesos

Turn Your Data Center Into One
Giant Computer



Agenda

- **Use Cases**
- What is Mesos
- Architecture
- Mesos Usage

System Are Getting Complicated

- More and tools/frameworks
 - Nodejs, J2EE, Ruby on Rails, Django, Flask, etc
 - MySQL, HBase, Cassandra, etc
 - Zookeeper, Etcd, etc
 - Hadoop, Spark, Samza, etc
- Mix of workloads
 - Long running services
 - Batch processes




Current Deployment Model

- Static cluster/data center partition
 - group similar apps into one server
- High operation code
- Vendor lock-in




Low Resource Utilization



- Different tasks generate different loads on machine
- Different tasks have different resource requirements
 - Memory bound - in memory data processing
 - CPU bound - rendering
 - IO bound - network apps



Management Layer for Data Center



Management Layer for Data Center





Agenda

- Use Cases
- **What is Mesos**
- Architecture
- Mesos Usage

What is Mesos

- An open-source cluster manager that
 - Turn datacenter/cluster into one computer
 - Provide simple API
 - Hide internal complex infrastructure from applications
- By Benjamin Hinderman, Andy Konwinski, and Matei Zaharia in UC Berkeley
- Inspired by Google Borg, matured in Twitter
- Apache Foundation top level project





Agenda

- Use Cases
- What is Mesos
- **Architecture**
- Mesos Usage



Architecture

- **Design Strategy**
- Concepts
- Internal

Two-level Scheduling

- Goal of Mesos is to schedule everything
 - Spark jobs
 - Jenkins build jobs
- Frameworks are different
 - API
 - Lifecycle
 - Scheduling requirements
- Best thing to do is to do nothing
 - Highly scalable
 - Small codebase
 - Easy to customize



Fairness

- How to fairly match jobs/tasks with resources
- How to work with different resource types
 - Memory
 - CPU





Architecture

- Design Strategy
- **Concepts**
- Internal

Framework

- Framework in Mesos is a distributed application developed with Mesos API
- A framework handles a type of workload or job
 - Hadoop
 - Spark
 - Jenkins
 - Chronos
 - Marathon



Resources and Attributes

- Resources represent what a slave has to offer
 - cpu
 - mem
 - disk
 - ports
 - `--resources='cpus:24;mem:24576;disk:409600;ports:[21000-24000];types:{a,b,c}'`
- Attributes are kv pairs that Mesos passes along when sending offer to framework
 - `--attributes='rack:abc;zone:west;os:centos5;level:10;keys:[1000-15000]'`



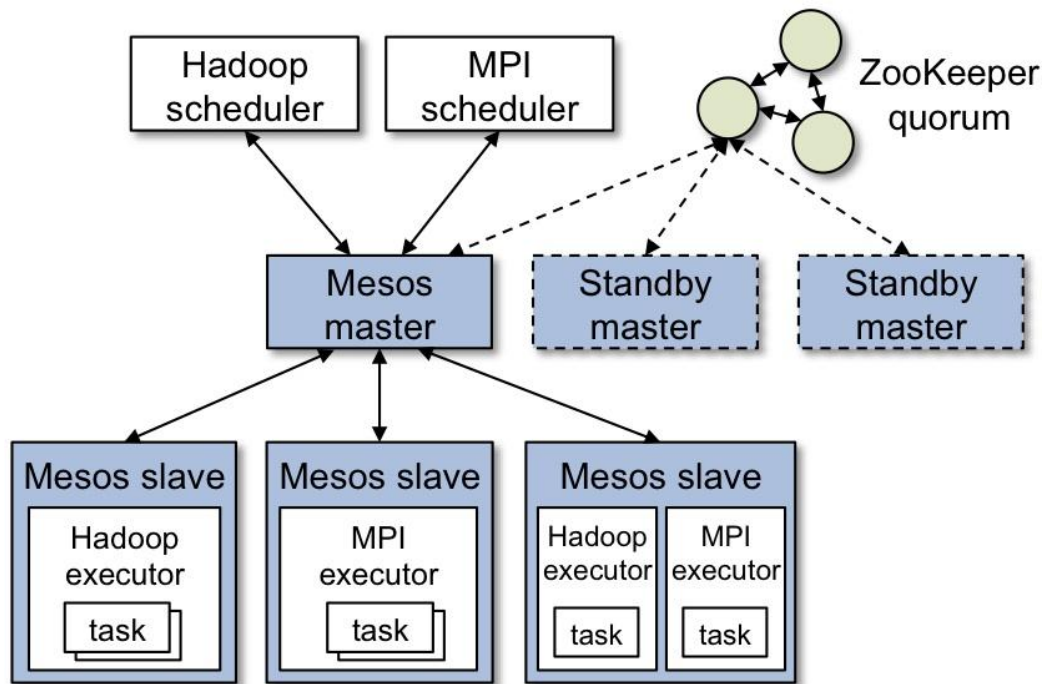


Architecture

- Design Strategy
- Concepts
- **Internal**

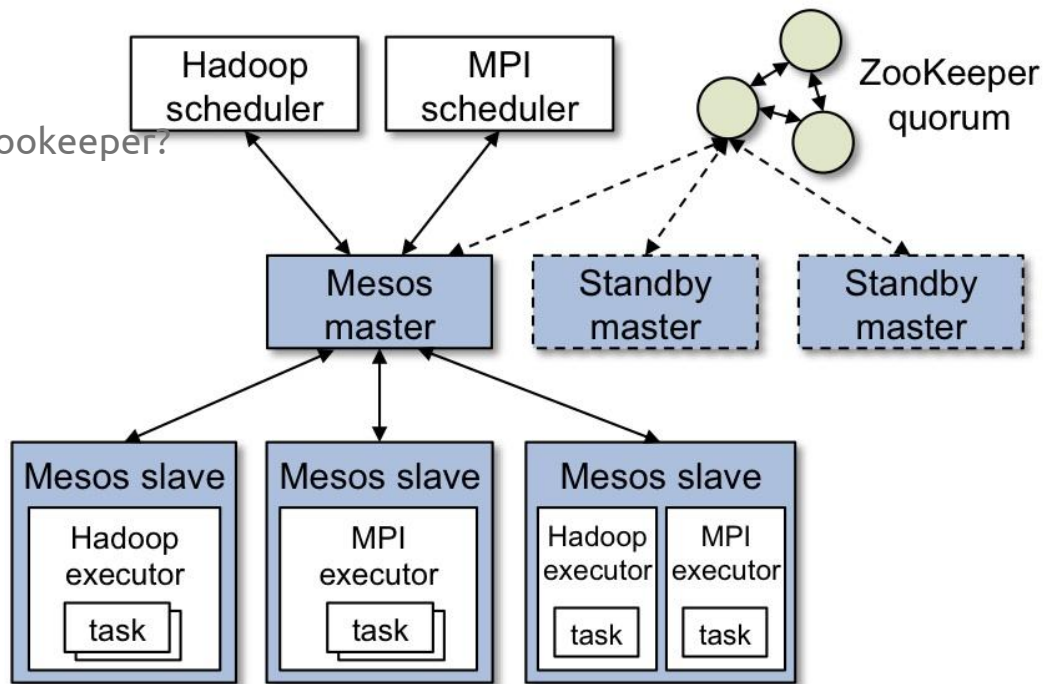
Main Components

- Mesos Master
- Mesos Slave
- Zookeeper
- Frameworks
 - Scheduler + Executor
- Protocol Buffer + libprocess



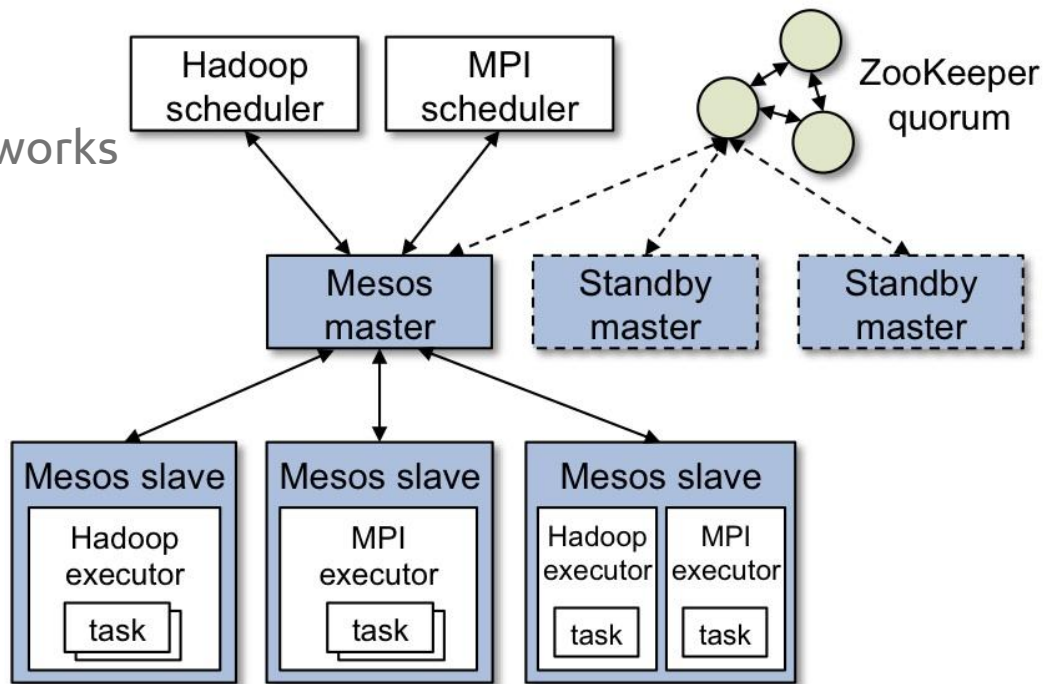
Zookeeper

- Handles leader election
 - How to do leader election with Zookeeper?




Mesos Master

- Manage Mesos slaves
- Make resource offer to frameworks



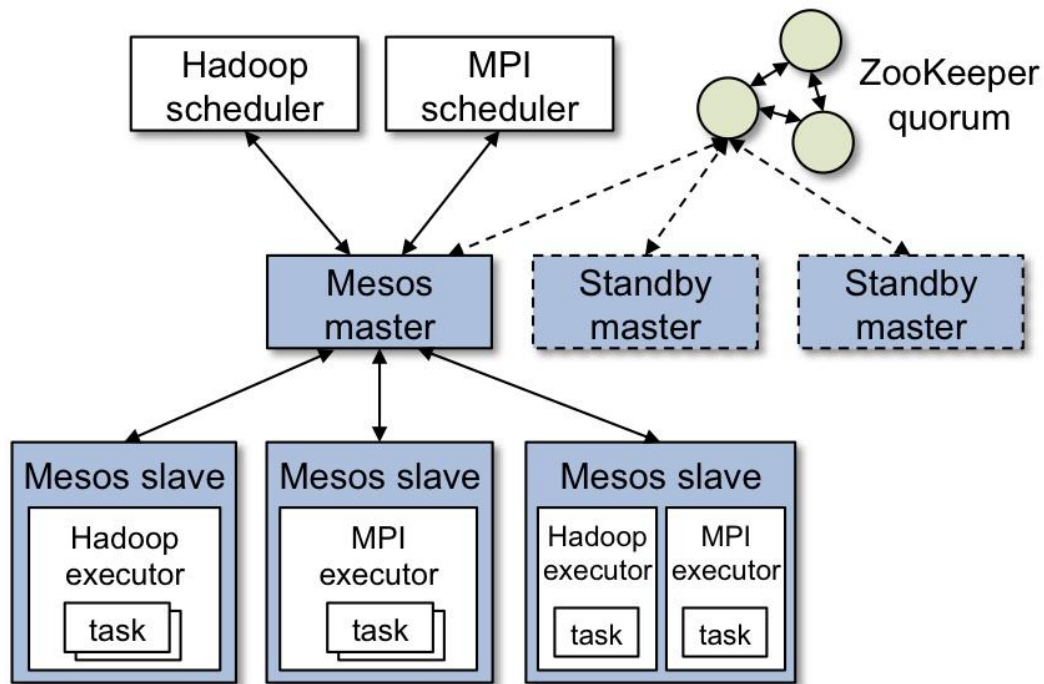
Mesos Slave

- Run tasks
- Report available resources
- Report task status, etc



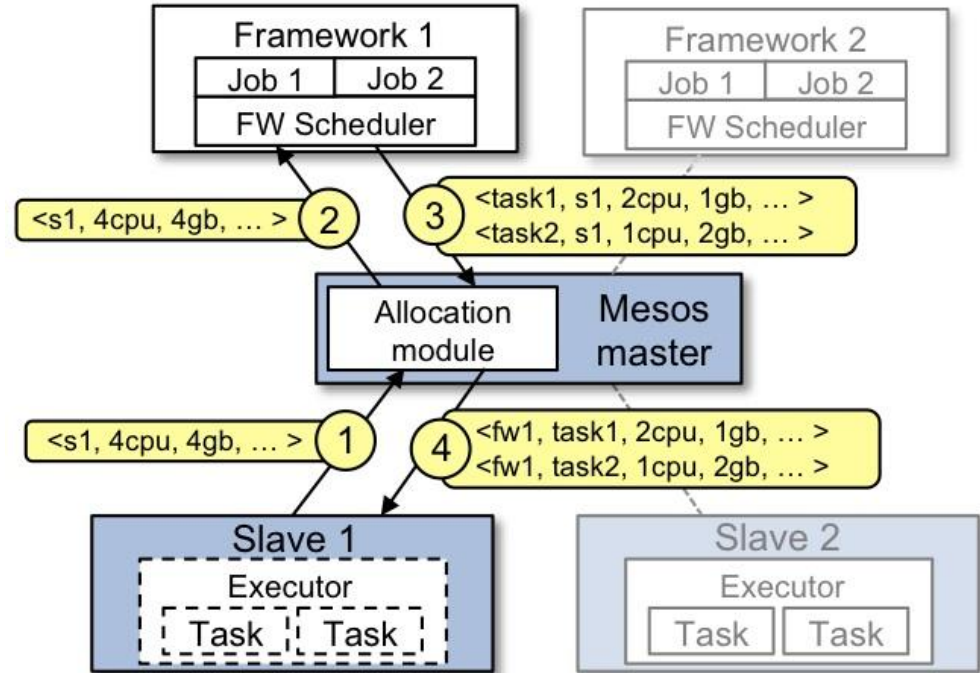
Mesos Framework

- Distributed Applications
- Scheduler processes offers and further schedule to tasks
- Executor launches tasks



Two Level Scheduling

- Allocation Module decide resources for each framework
- Framework Scheduler decide resources for each tasks



Scheduling Algorithm

- Dominant Resource Fairness Algorithm (DRF)
- Based on Benjamin Hinderman's paper
- Default Mesos Allocation Module implementation
- Achieves fairness (maybe too much?)



An Example Problem

- Consider we have one pizza, 10 slices
- A want 1 slice, B want 2 slices, C want 5 slices, and D want 6 slices.
- How to achieve fairness in this case?
 - Divide equally - everyone get 2.5 slices, but C will get hungry
 - What if someone lied about their resource request?



Max-min Fairness Algorithm

- Satisfy small tasks first
- Divide the rest resources among others
 - A wants 1 slice, get 1 slice, 9 slices left
 - B wants 2 slices, get 2 slices, 7 slices left
 - C wants 5 slices, not enough left, divide with other people, get 3.5 slices, 3.5 slices left
 - D get 3.5 slices
- Max-min Fairness is a good solution for fairness



Dominant Resource Fairness (DRF)

- DRF is a max-min fair algorithm for heterogeneous resources
 - CPU
 - Memory
 - IO



Dominant Resource Fairness (DRF)

- Nice qualities for scheduling algorithms
 - If C gets less than 2.5 slices, he will just go and buy his own pizza - Sharing Incentive
 - B lies about his share, what would happen? - Strategy Proofness
 - C should not envy B share - Envy Freeness
 - You cannot give B more pizza without give other people less pizza - Pareto Efficiency



Dominant Resource Fairness (DRF)

- Dominant Resource Share
 - If we have a system of 10 CPU 10 GB RAM, User A has been assigned with 2 CPU and 6 GB RAM, Dominant Resource Share is $\frac{3}{5}$, and A's dominant resource is memory.
- DRF tries to maximize min dominant resource share.

	Framework 1			Framework 2			CPU	RAM
Framework Chosen	Resoure Shares	Dominant Share	Dominant Share %	Resoure Shares	Dominant Share	Dominant Share %	Total Allocation	Total Allocation
	0/9, 0/18	0	0%	0/9, 0/18	0	0%	0/9	0/18
Framework 2	0/9, 0/18	0	0%	3/9, 1/18	1/3	33%	3/9	1/18
Framework 1	1/9, 4/18	2/9	22%	3/9, 1/18	1/3	33%	4/9	5/18
Framework 1	2/9, 8/18	4/9	44%	3/9, 1/18	1/3	33%	5/9	9/18
Framework 2	2/9, 8/18	4/9	44%	6/9, 2/18	2/3	67%	8/9	10/18
Framework 1	3/9, 12/18	2/3	67%	6/9, 2/18	2/3	67%	9/9	14/18





Agenda

- Use Cases
- What is Mesos
- Architecture
- **Mesos Usage**



Mesos Usage

- Shared Resource Pool
- Container Orchestration

Shared Resource Pool

- Run mix of short tasks on Mesos
 - Jenkins - build jobs, deploy jobs
 - Spark
 - Hadoop
 - Ad-hoc queries
- Apple runs Siri on top of Mesos




Container Orchestration

- Use Apache Marathon (<https://github.com/mesosphere/marathon>)
 - High availability, failed containers will be respawned
 - Automatically load balancing
 - Service discovery



Container Orchestration



Applications

Deployments

About

API Reference

Documentation

Create

Applications

Search all applications



STATUS

☐ Running 1

☐ Deploying

☐ Suspended 1

☐ Delayed

☐ Waiting

HEALTH

☐ Healthy 1

☐ Unhealthy

☐ Unknown

Name ▲

CPU

Memory

Status

Running Instances

Health



helloworldconnected

HAPROXY_0_VHOST:foo.com

0.1

64 MiB

Running

1 of 1

...



website

0.0

0 B

Suspended

0 of 0

...

Similar Systems

- YARN (<https://hortonworks.com/apache/yarn>)
 - Developed to improve Hadoop resource efficiency
- Kubernetes (<http://kubernetes.io/>)
 - Container orchestration solution from Google
 - Based on Borg targeting Omega
- Docker Swarm (<https://docs.docker.com/swarm/>)
 - Native Container orchestration solution from Docker



Further Reading

- Borg: <http://research.google.com/pubs/pub41684.html>
- Borg, Omega, and Kubernetes:
<http://research.google.com/pubs/pub44843.html>
- Omega: <http://research.google.com/pubs/pub43438.html>
- Mesos: <https://www.cs.berkeley.edu/~alig/papers/mesos.pdf>
- DRF: <https://www.cs.berkeley.edu/~alig/papers/drf.pdf>
- Orchestration: [https://en.wikipedia.org/wiki/Orchestration_\(computing\)](https://en.wikipedia.org/wiki/Orchestration_(computing))





Q&A